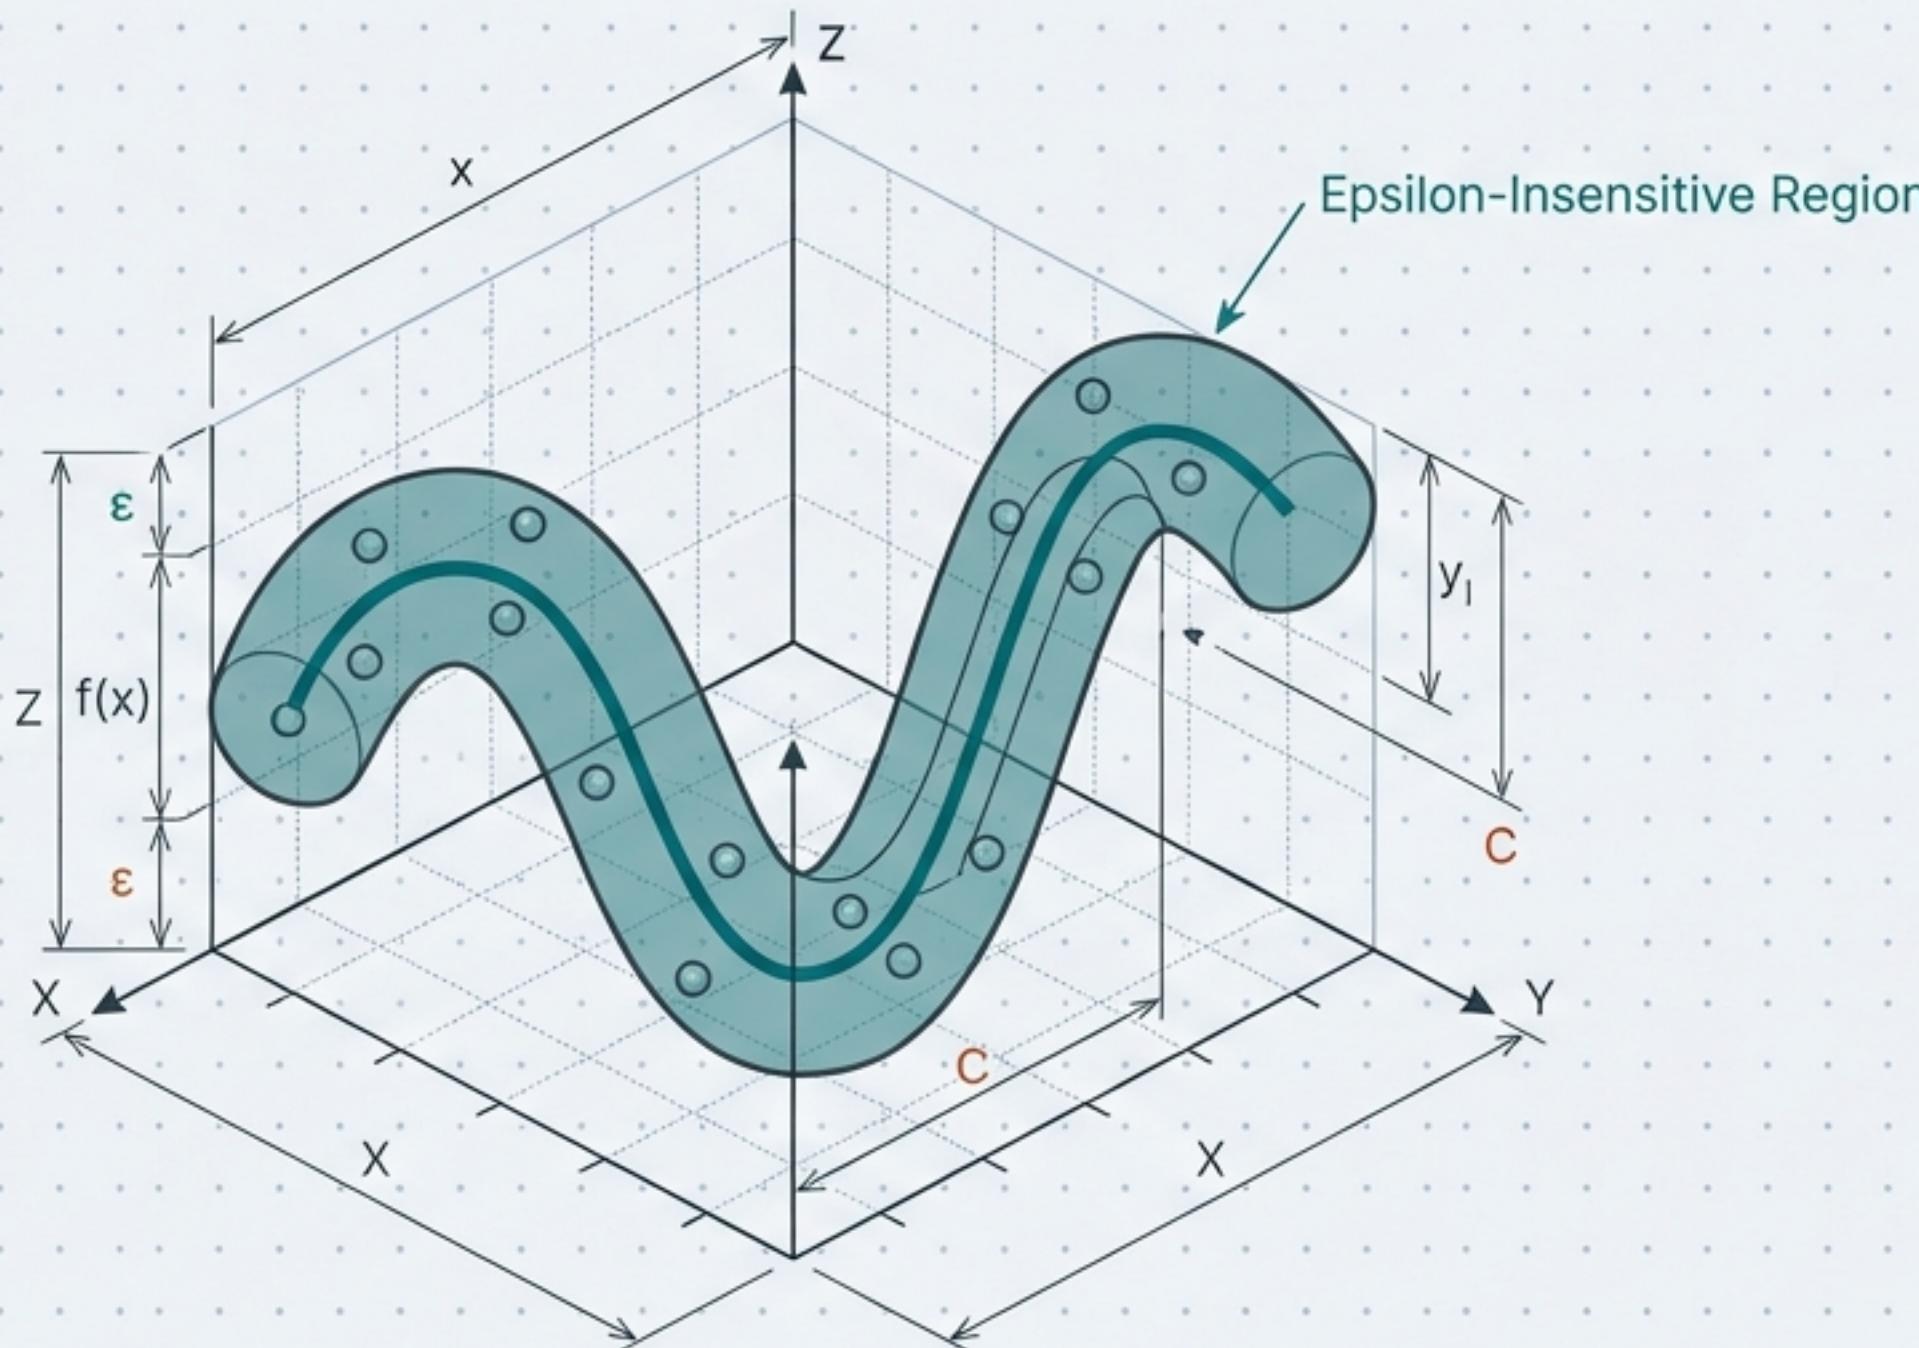


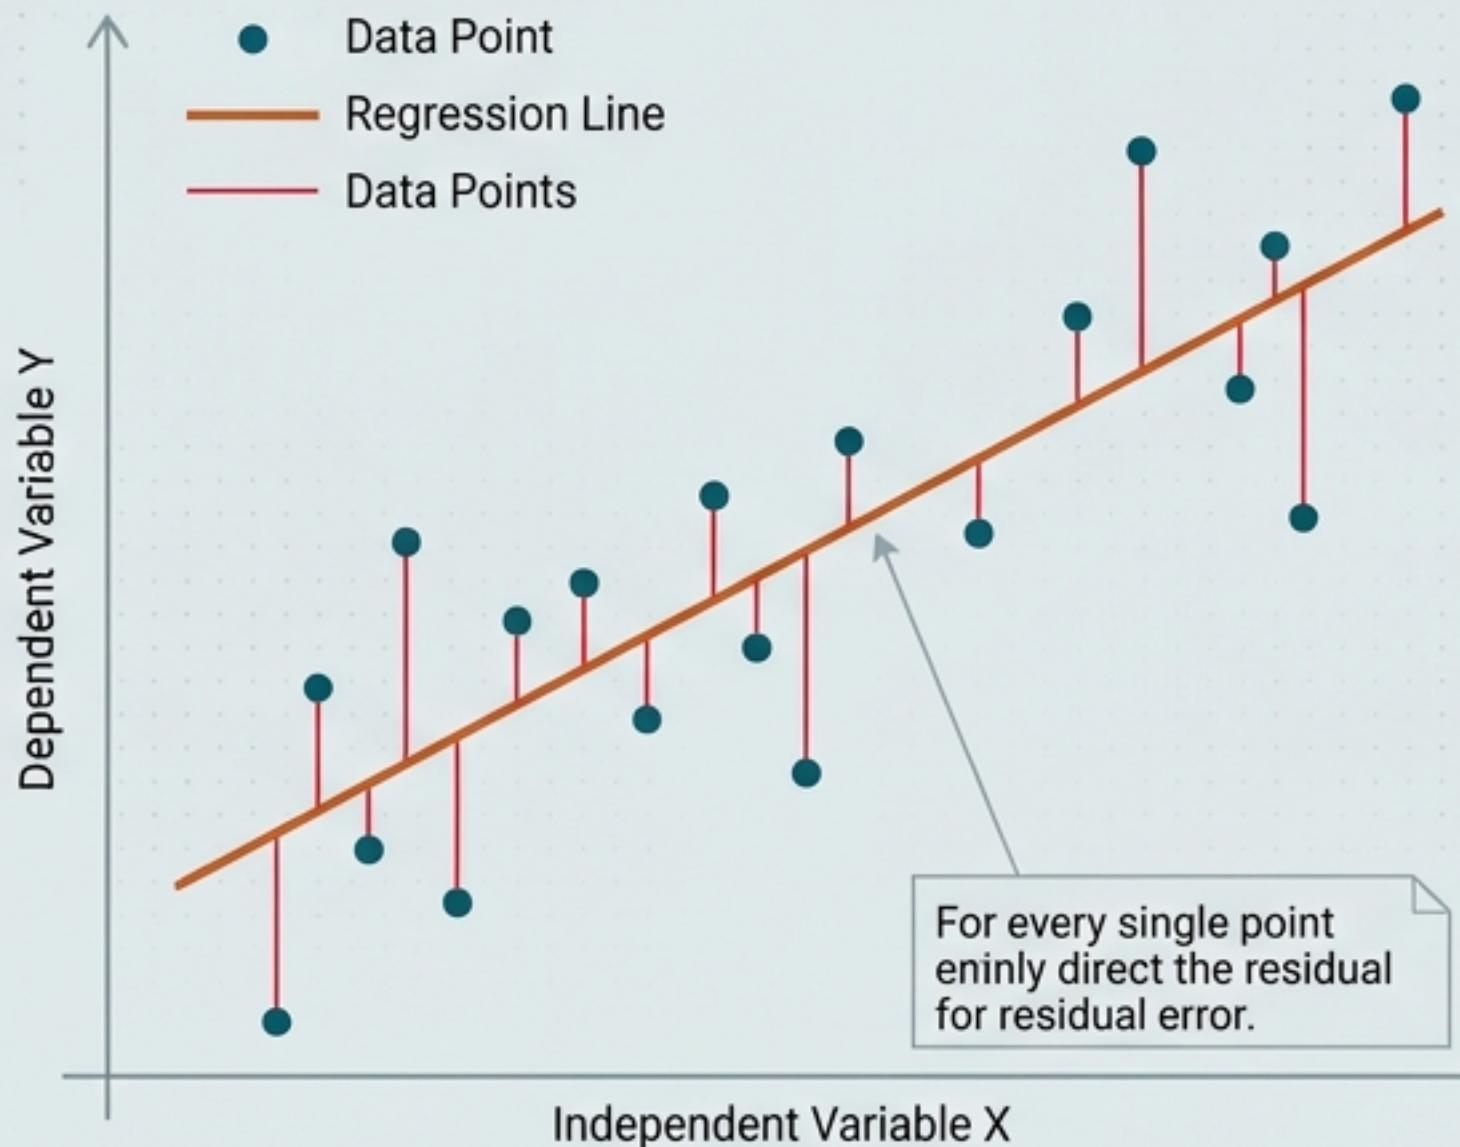
# Mastering Support Vector Regression (SVR)

A Guide to Predictive Modeling with Tolerance



# The Paradigm Shift: From Bullseye to Dartboard

## OLS (Ordinary Least Squares)



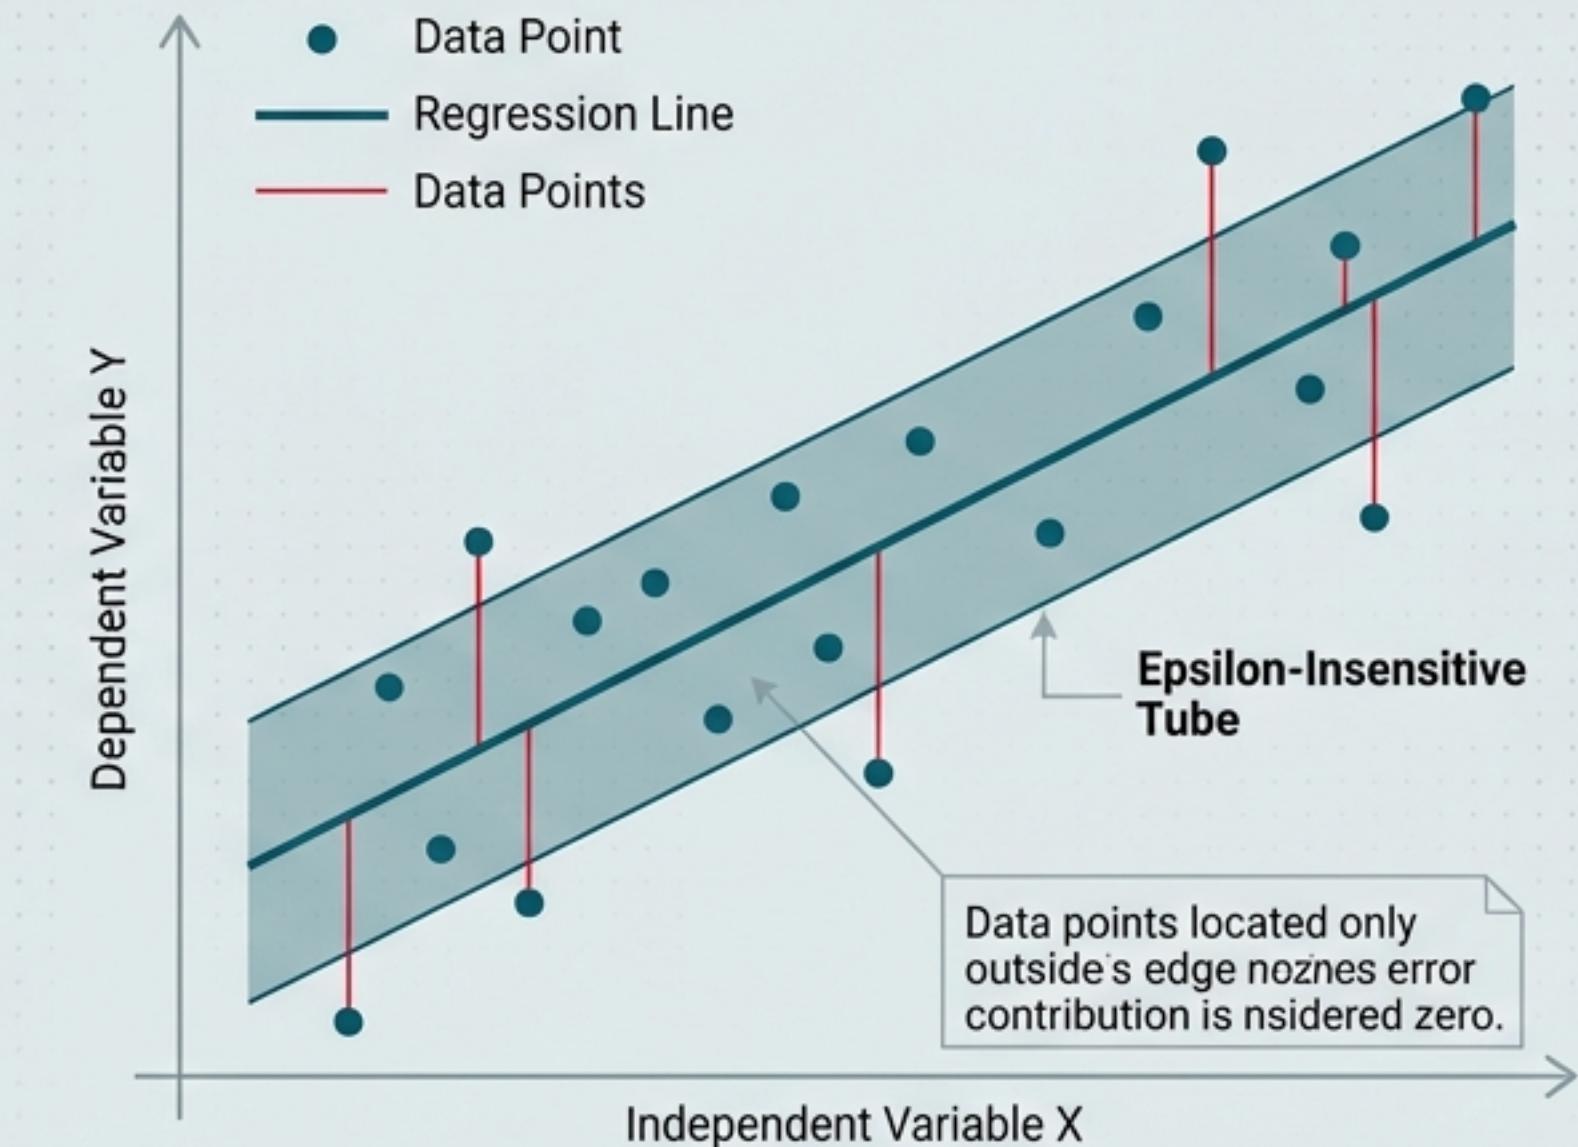
### Minimizing Sum of Squared Errors

OLS tries to hit the bullseye.

Every error counts, making it sensitive to outliers.

**Empirical Risk Minimization**

## SVR (Support Vector Regression)



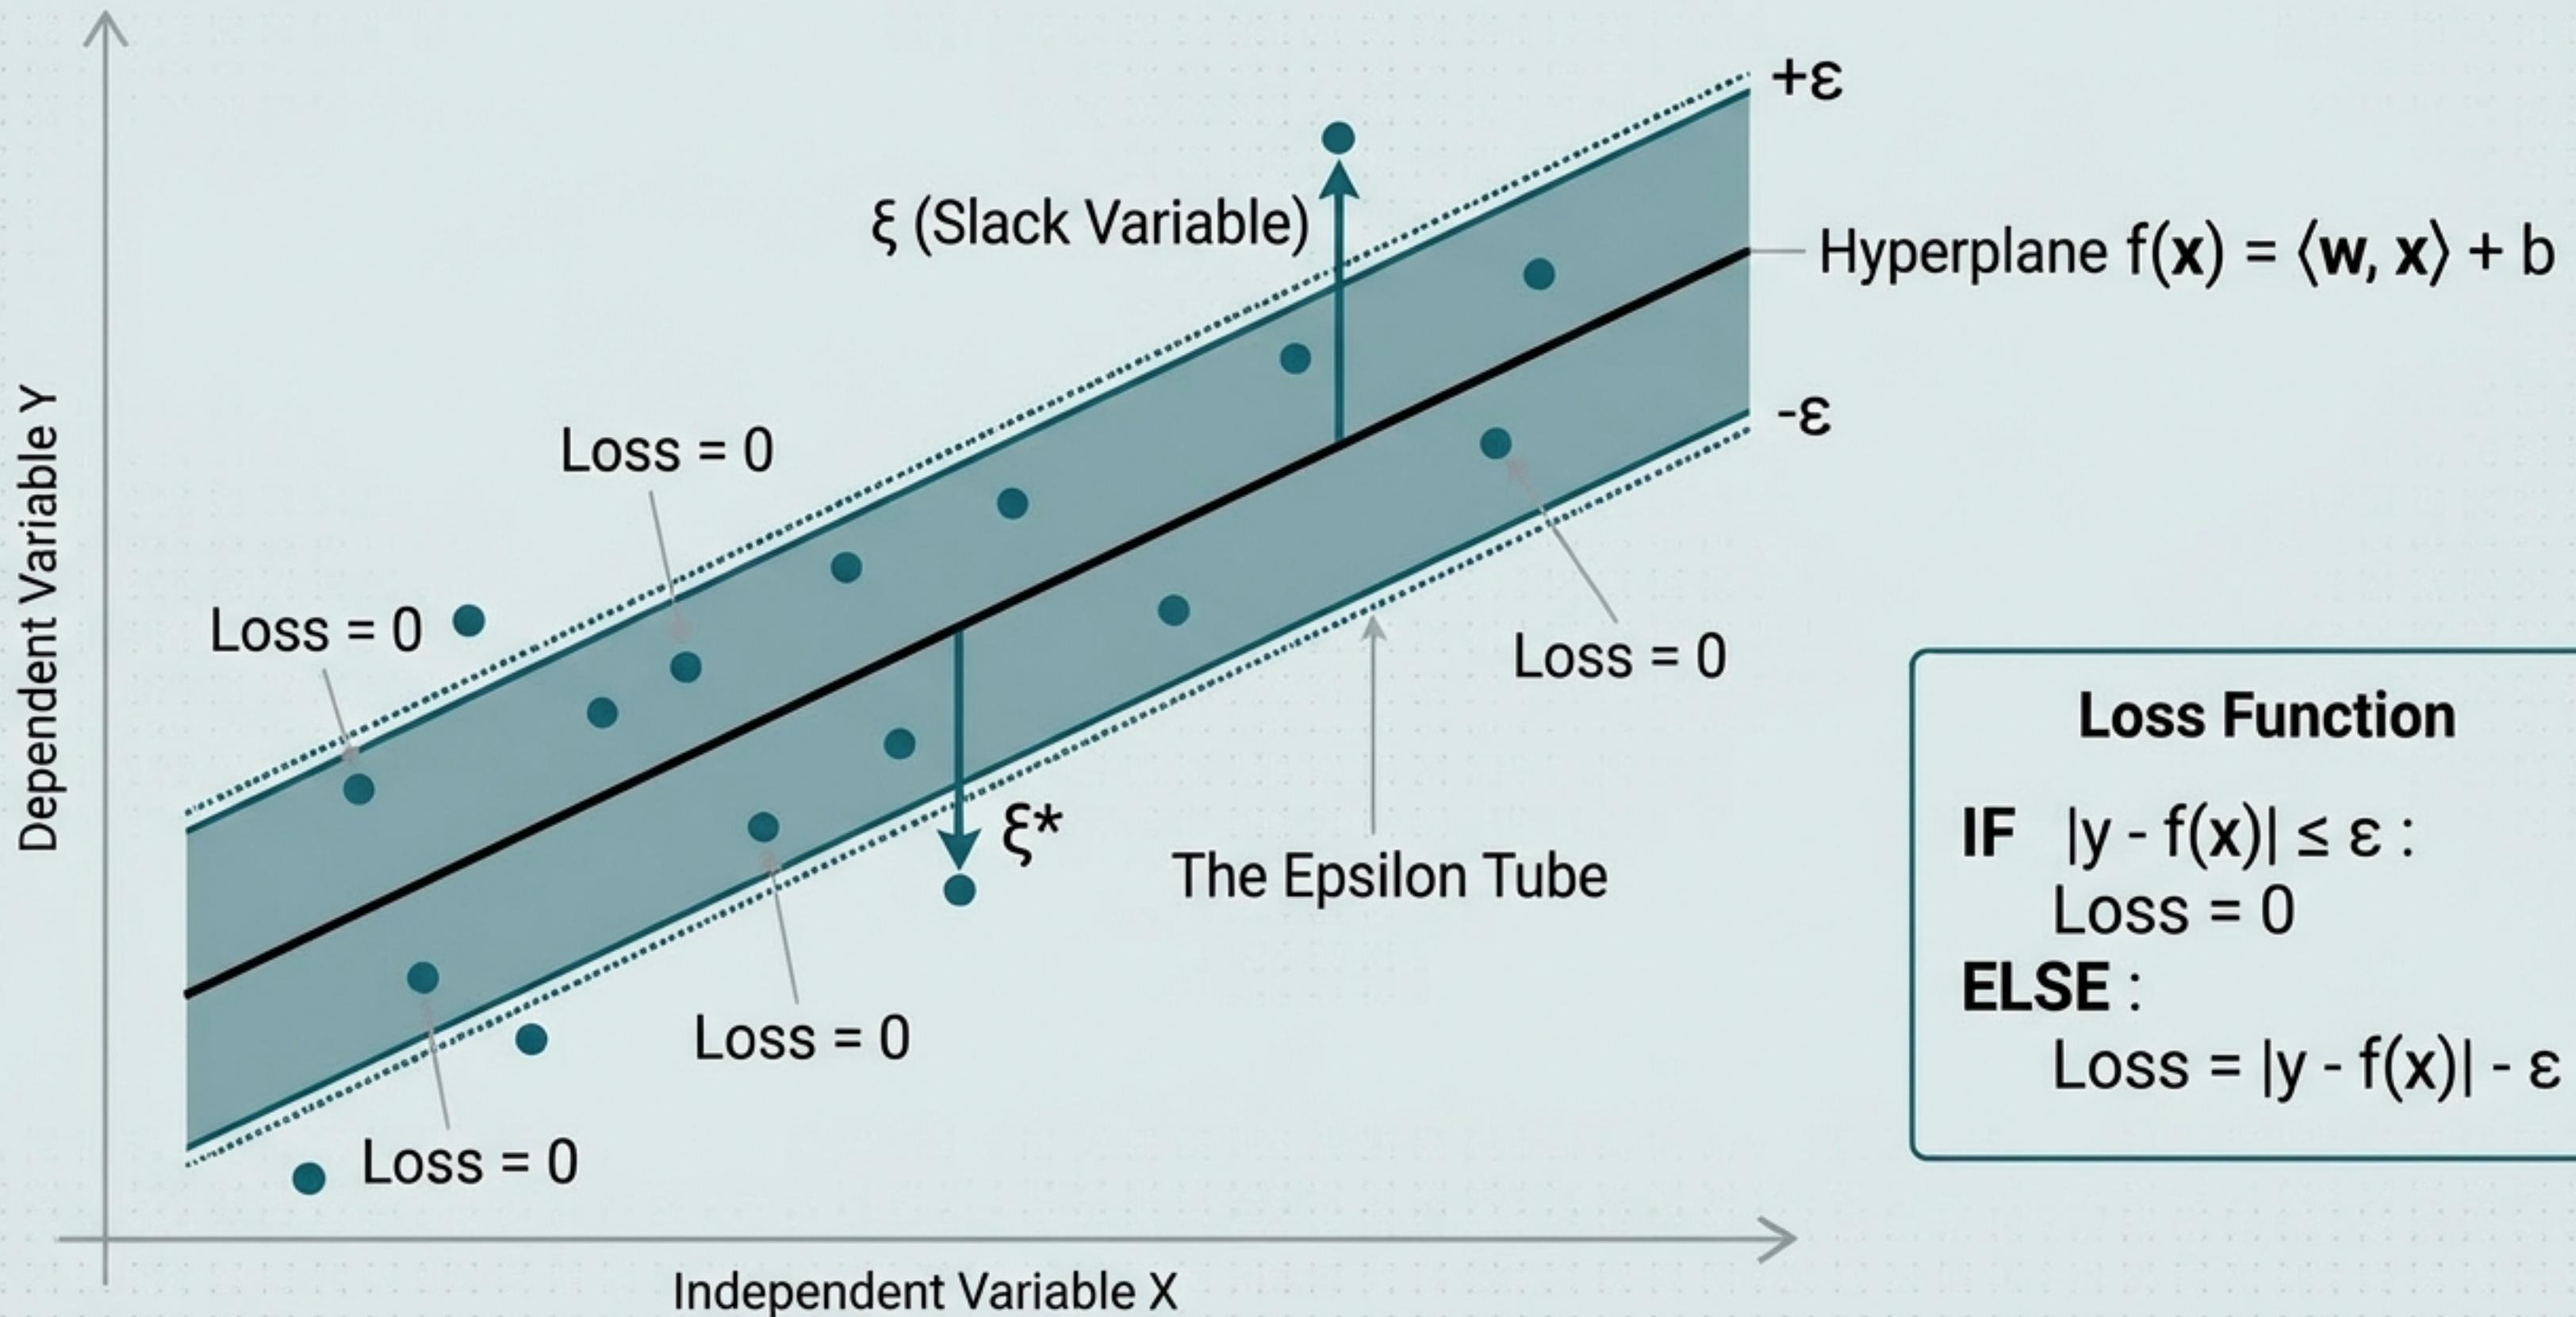
### Epsilon-Insensitive Loss

SVR tries to hit the dartboard.

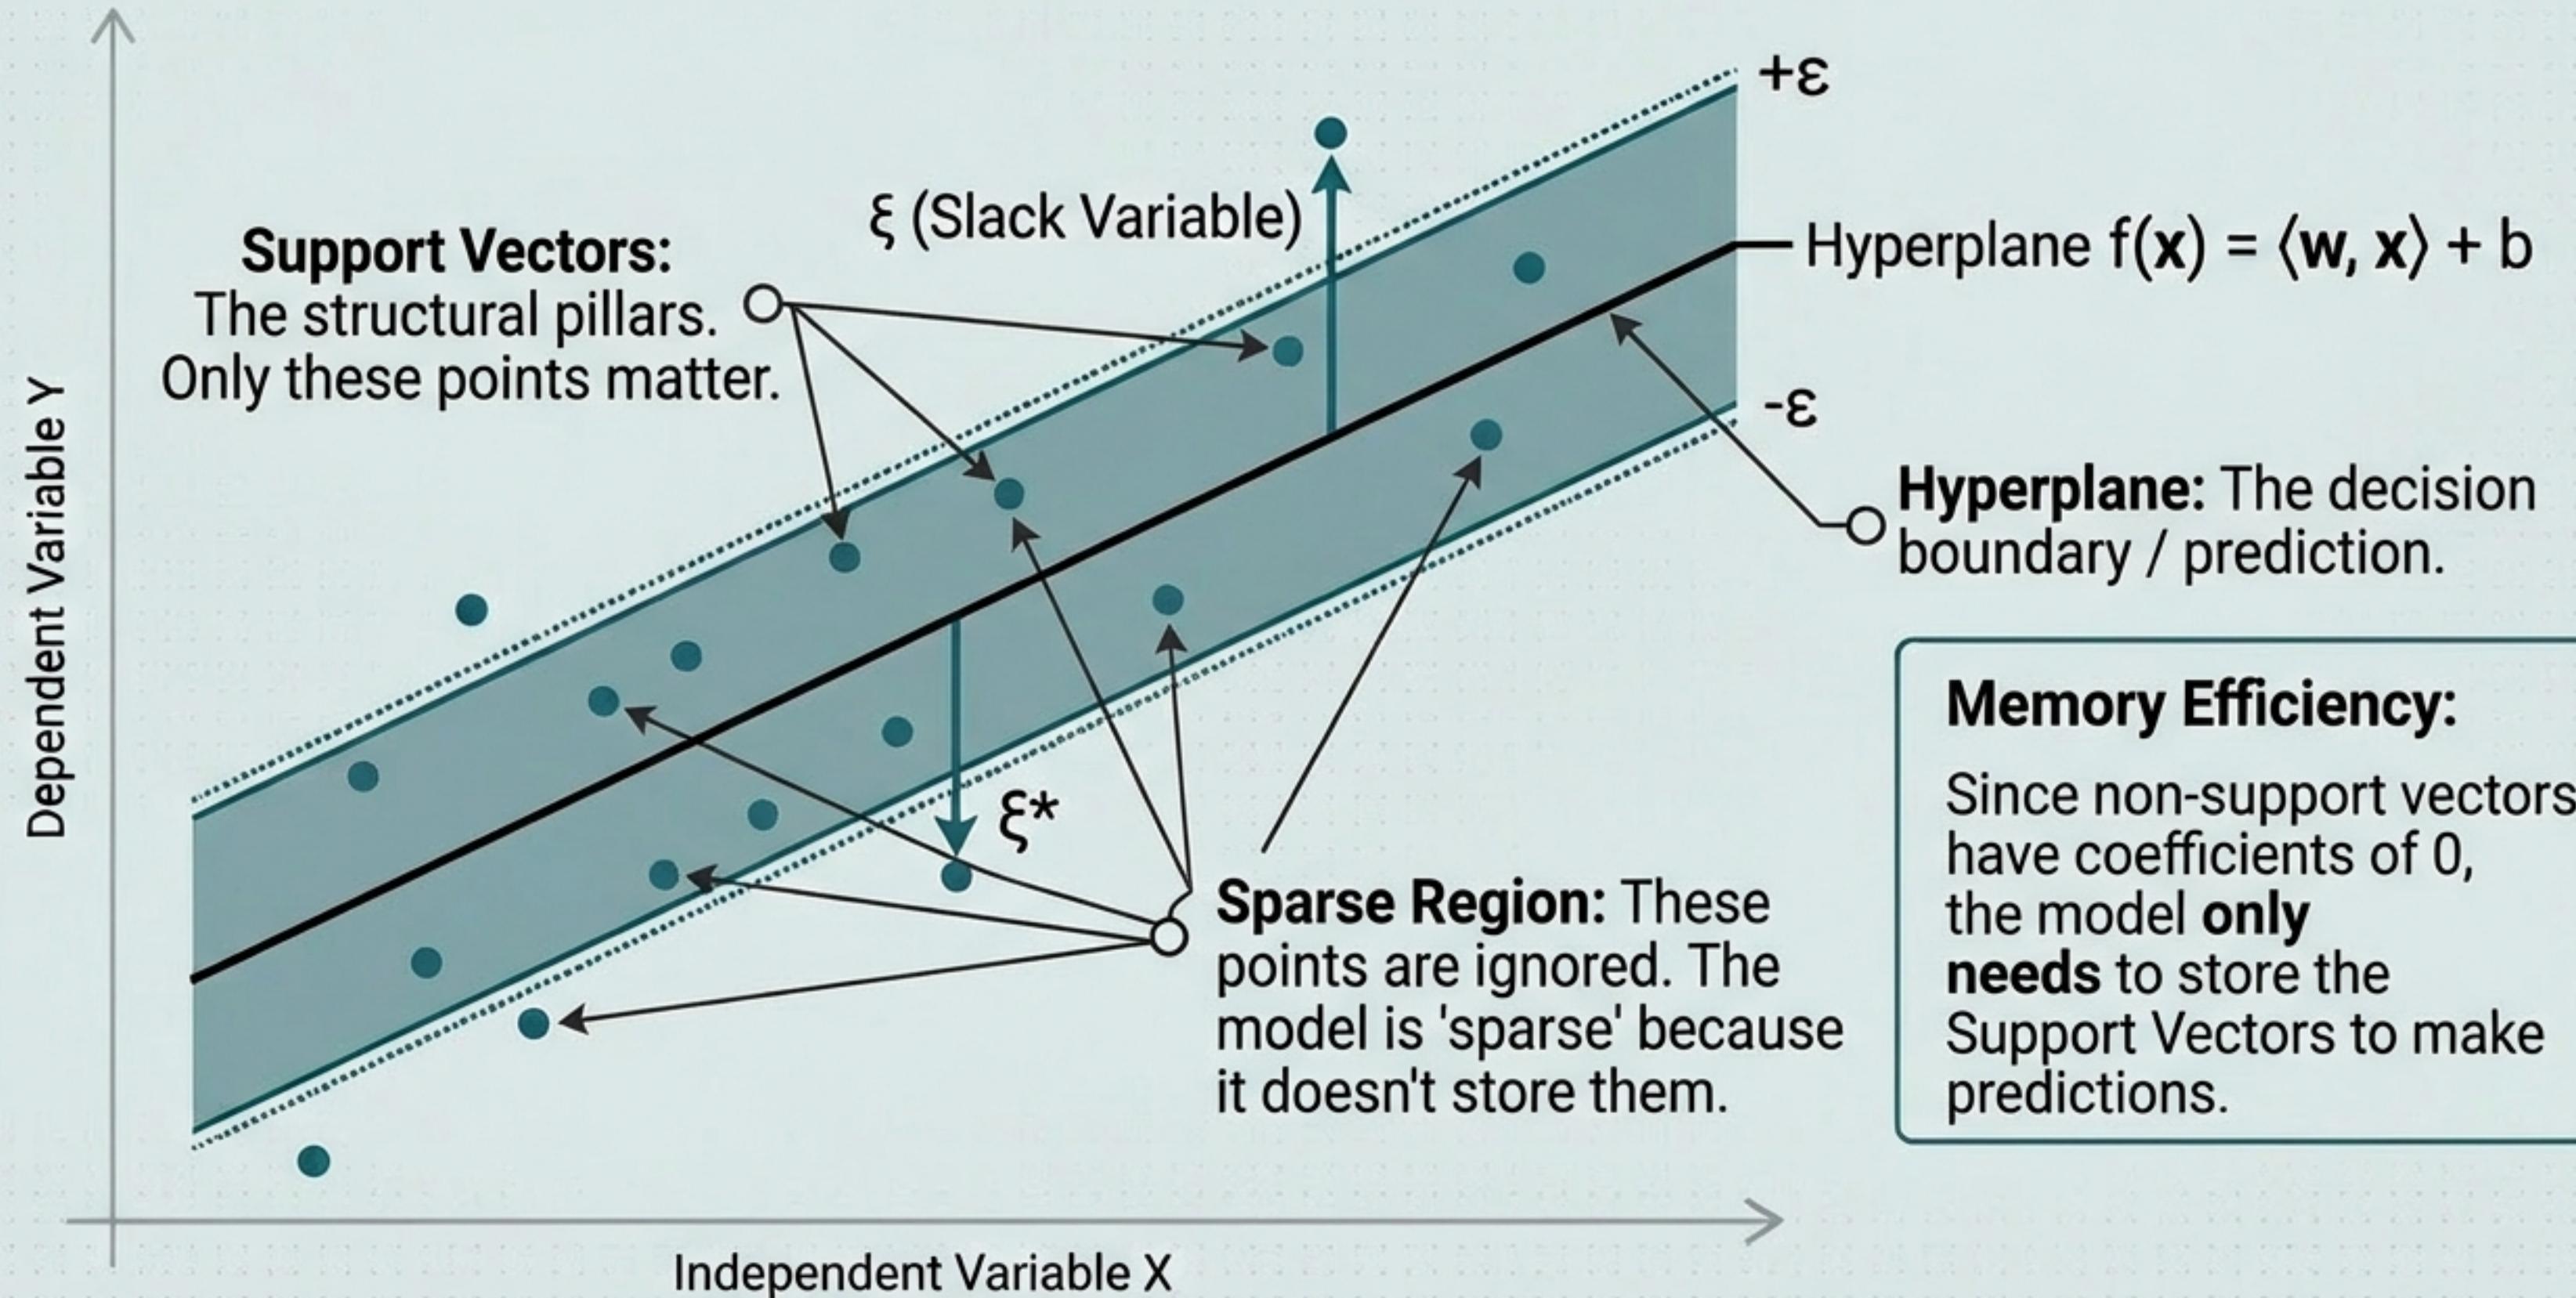
"Close enough" is good enough.  
Errors inside the tube are ignored.

**Structural Risk Minimization**

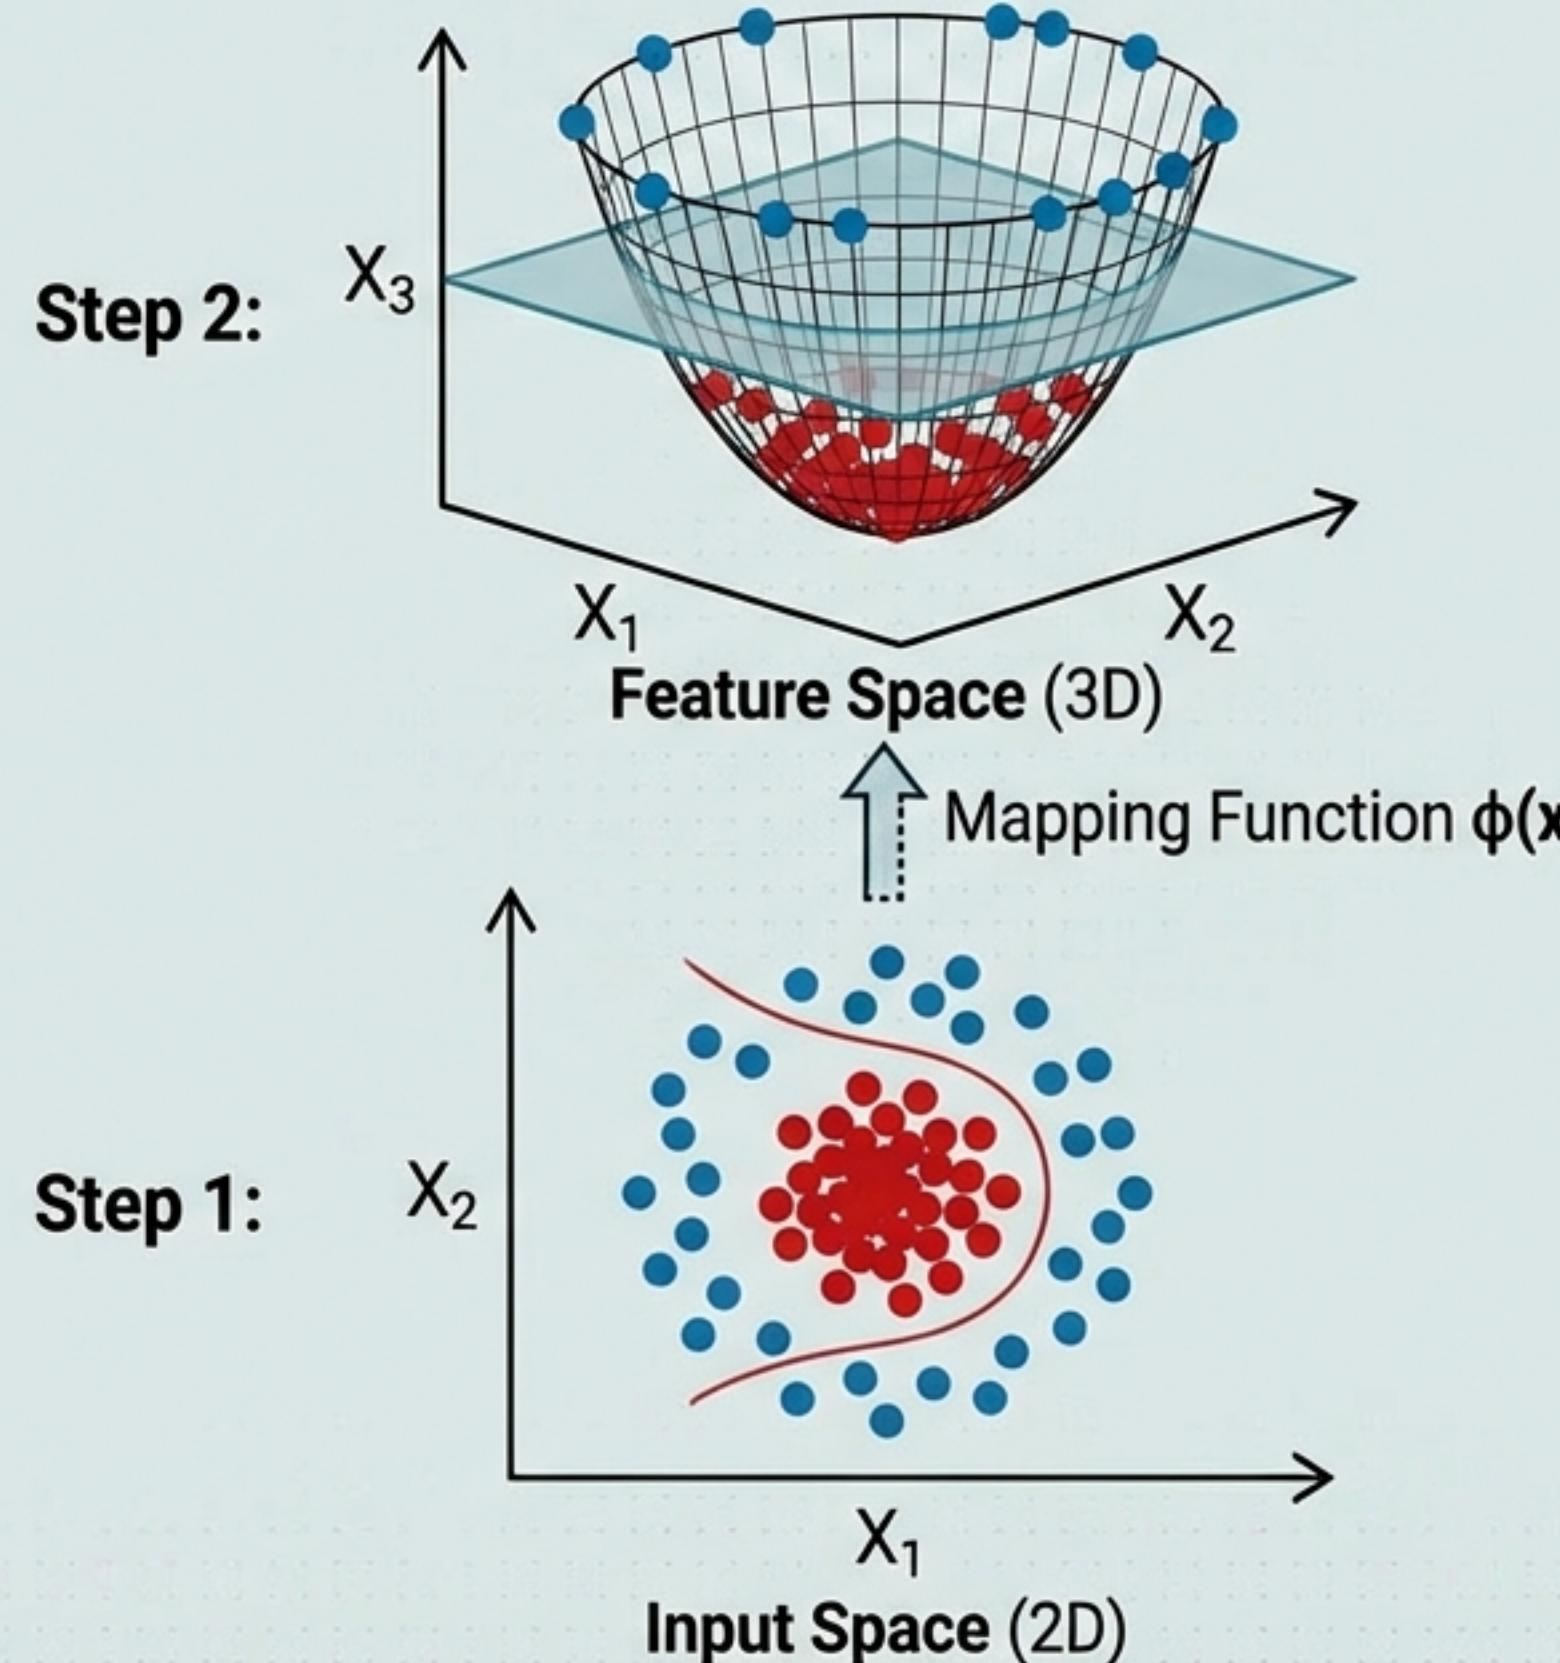
# The Geometry of Tolerance: The $\varepsilon$ -Insensitive Tube



# The Cast of Characters



# The Kernel Trick: Fitting Curves with Straight Sheets



**Linearly Separable.**

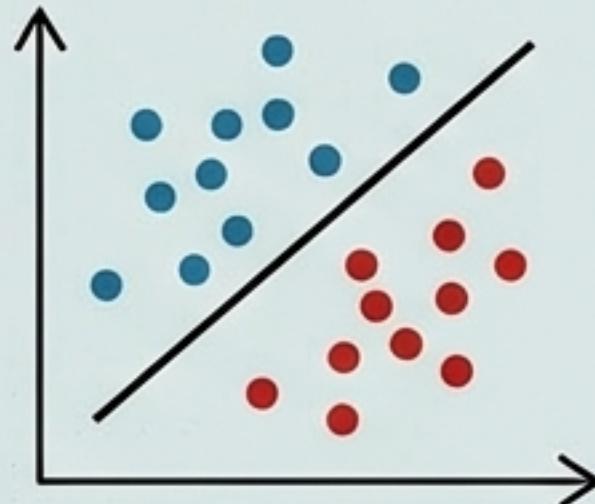
The kernel trick computes dot products in this space without explicitly calculating coordinates.

**Linearly Inseparable.**  
No straight line can separate red from blue.

# The Kernel Menu: Choosing Your Geometry

## Linear Kernel

`kernel='linear'`

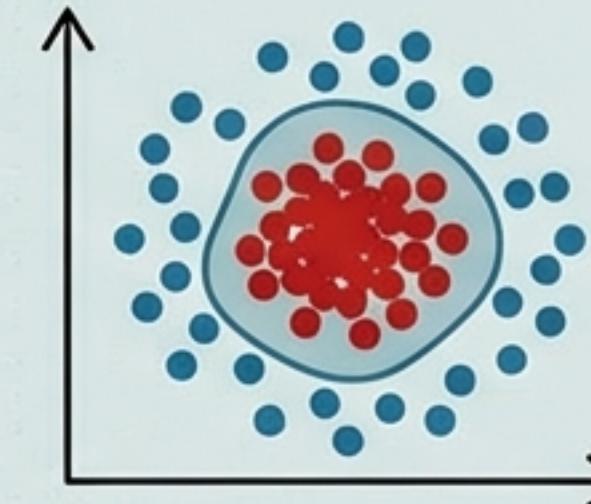


**Math:**  $\langle \mathbf{x}, \mathbf{x}' \rangle$ .

**Best for:** High-dimensional data (Text), massive datasets. Fast.

## RBF Kernel (Radial Basis Function)

`kernel='rbf'`

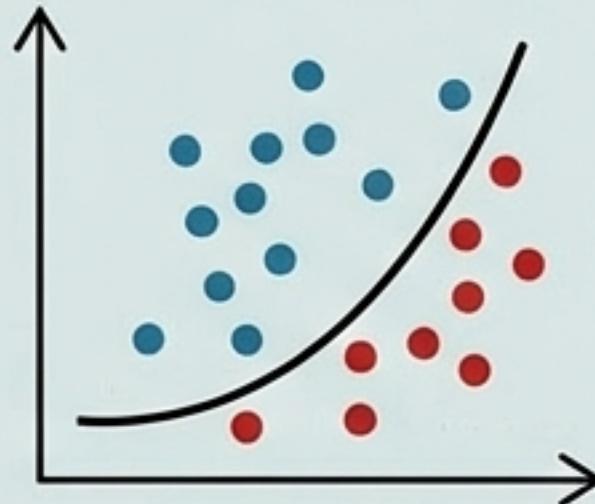


**Math:**  $\exp(-\gamma \|\mathbf{x} - \mathbf{x}'\|^2)$

**Best for:** The Default choice. Maps to infinite dimensions. Captures local behavior.

## Polynomial Kernel

`kernel='poly'`

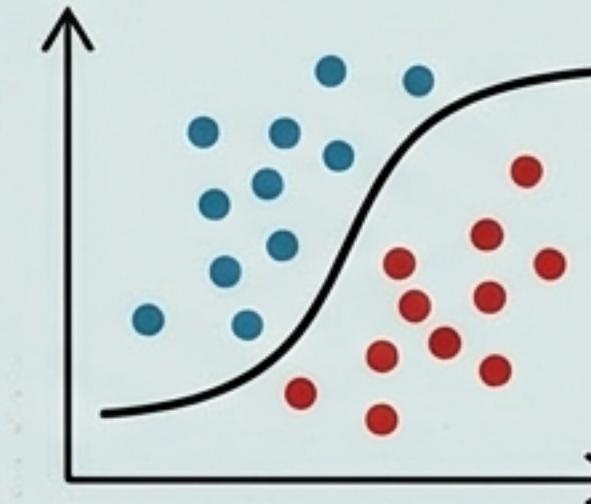


**Math:**  $(\gamma \langle \mathbf{x}, \mathbf{x}' \rangle + r)^d$

**Best for:** Image processing. Captures feature interactions.

## Sigmoid Kernel

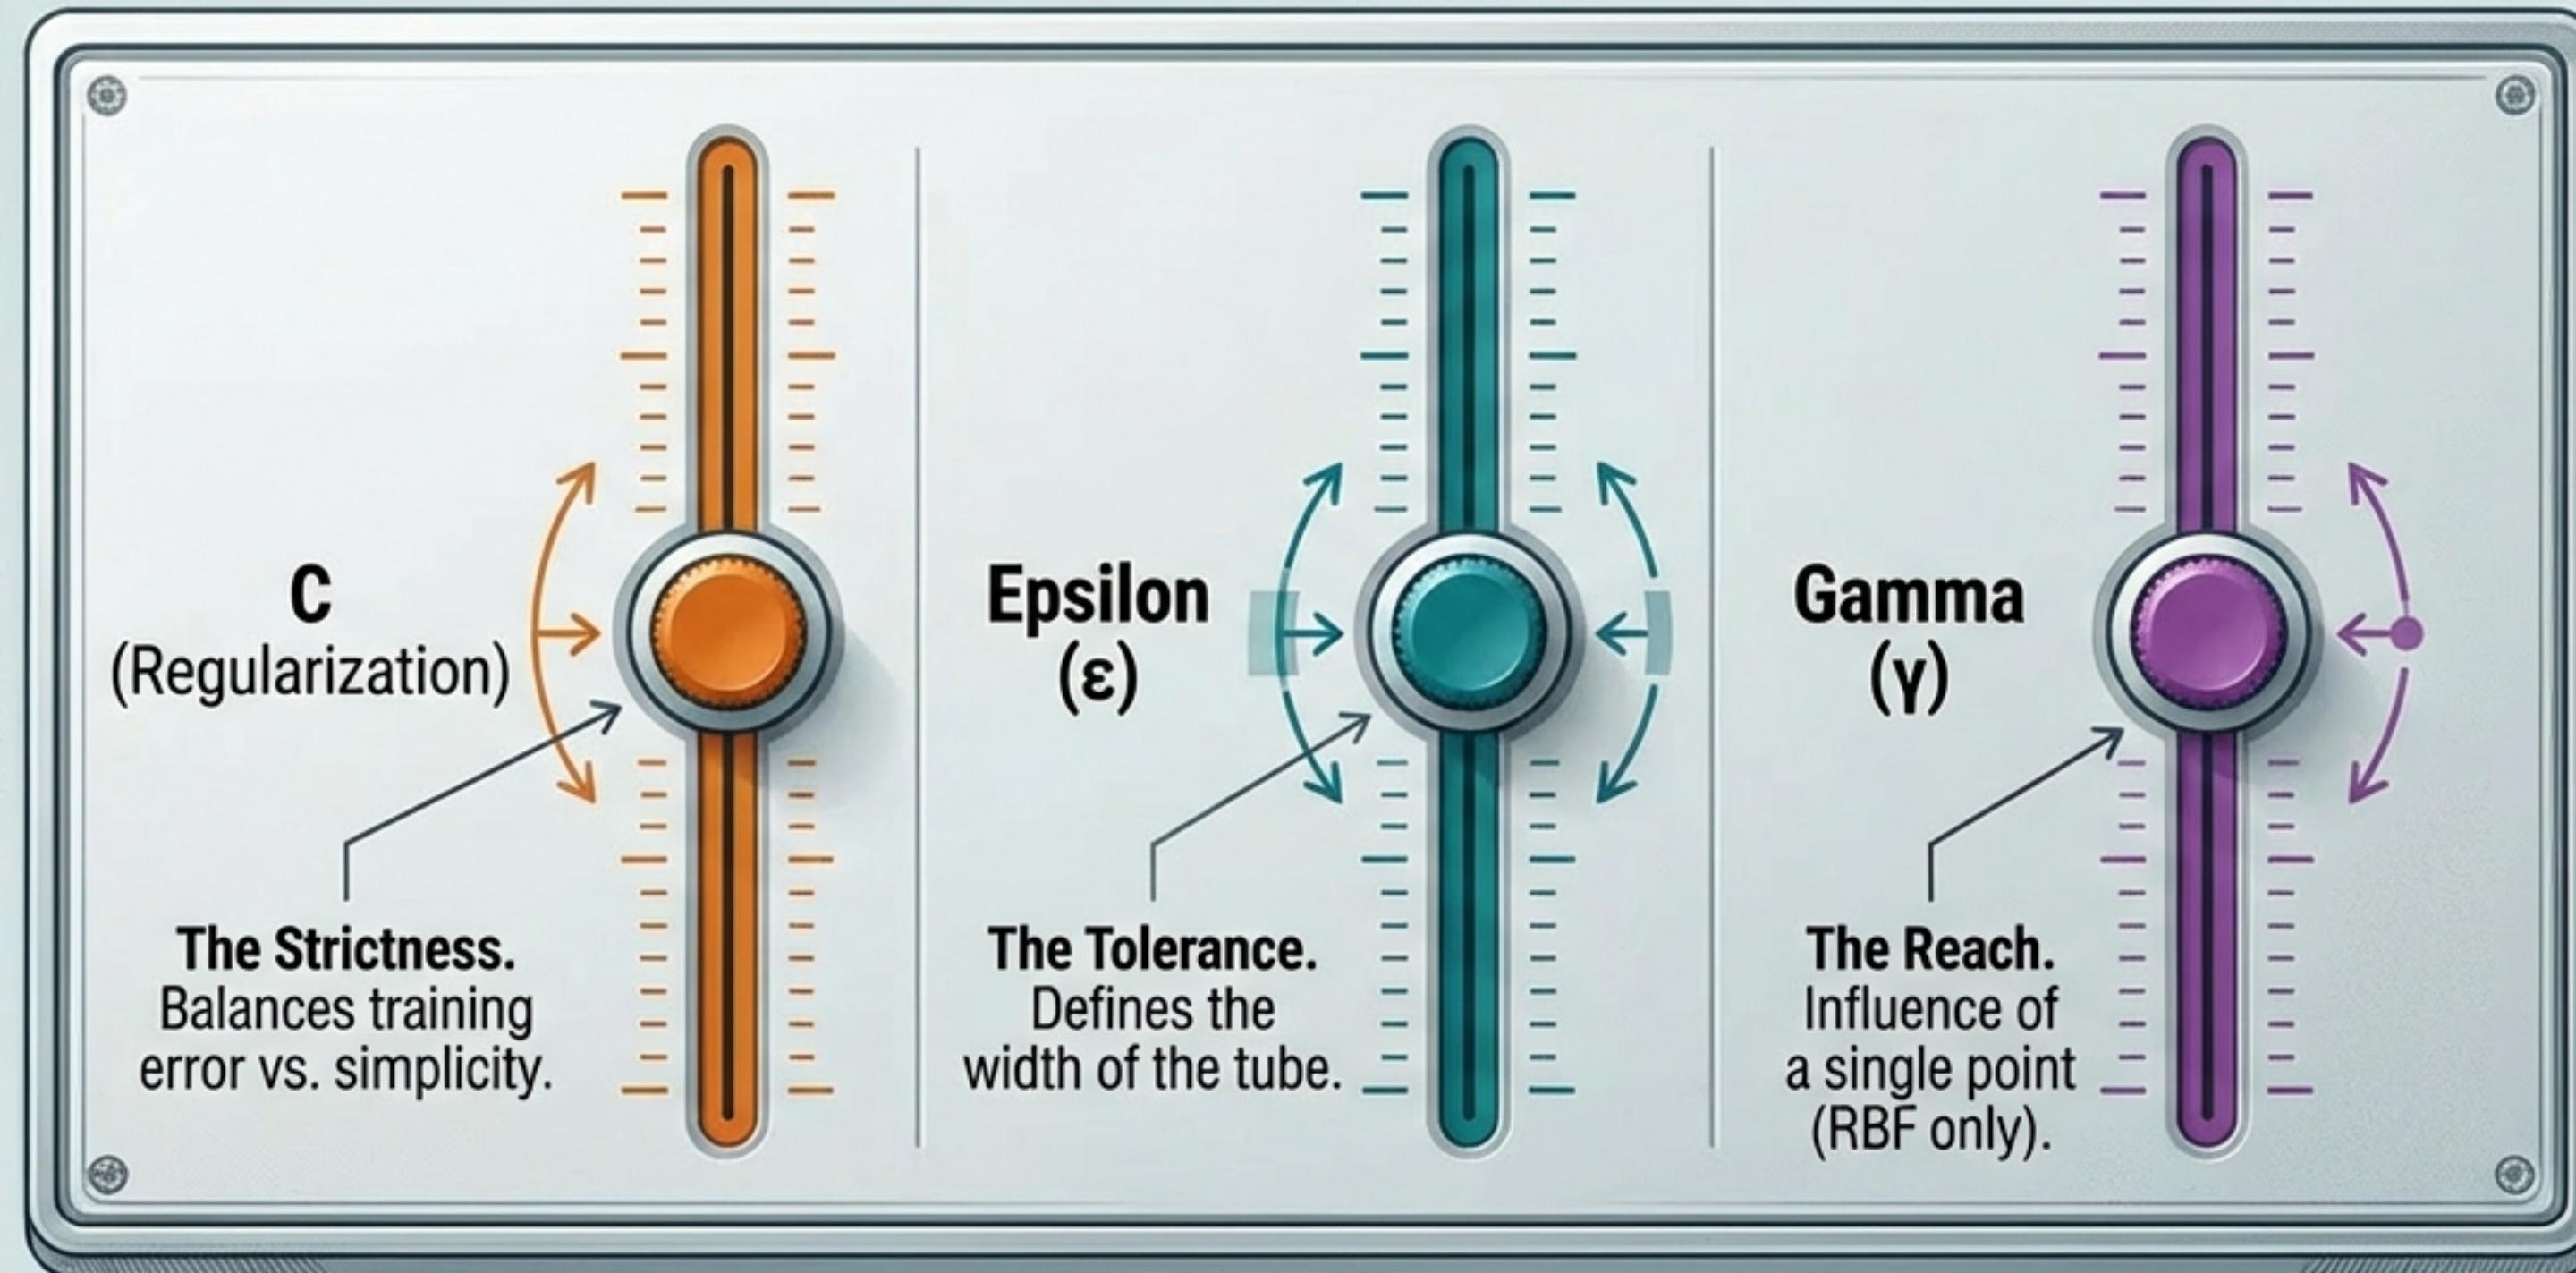
`kernel='sigmoid'`



**Math:**  $\tanh(\gamma \langle \mathbf{x}, \mathbf{x}' \rangle + r)$

**Best for:** Neural Network proxy. Often unstable for general regression.

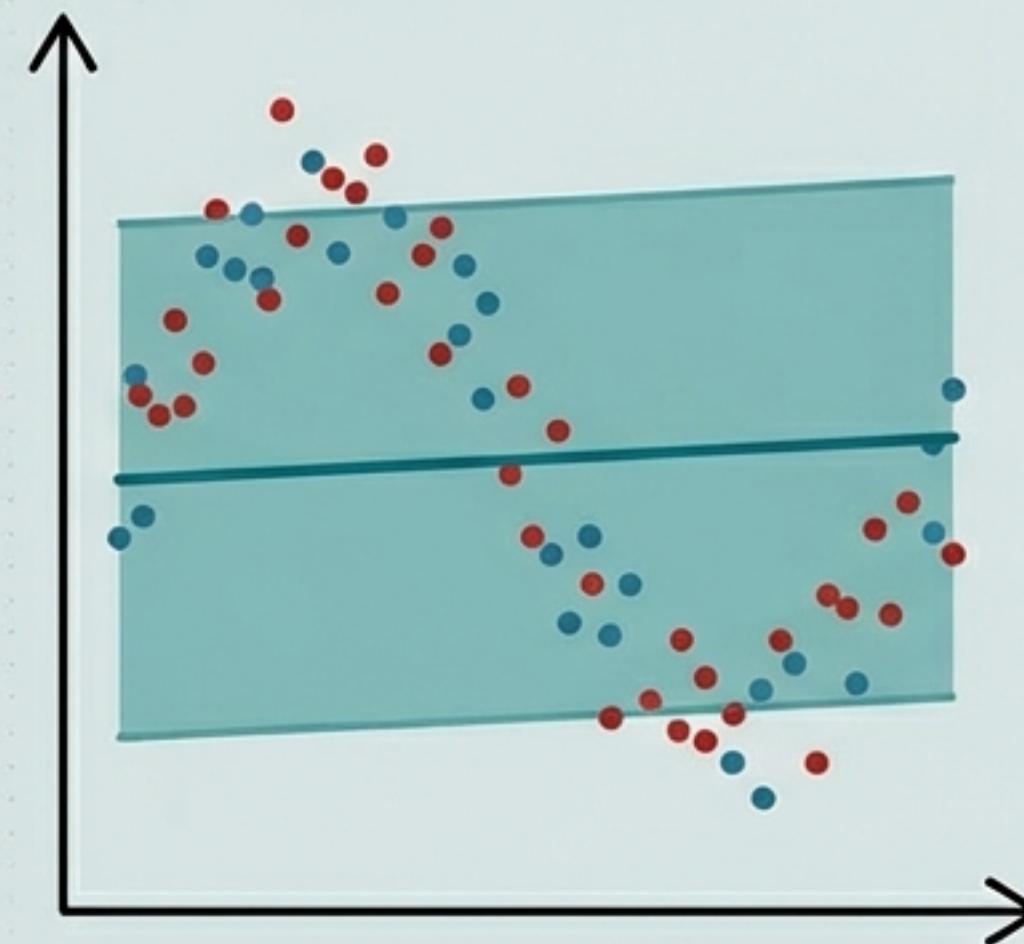
# The Hyperparameter Ecosystem



**Ecology:** These parameters are interconnected. Changing one often requires tuning the others.

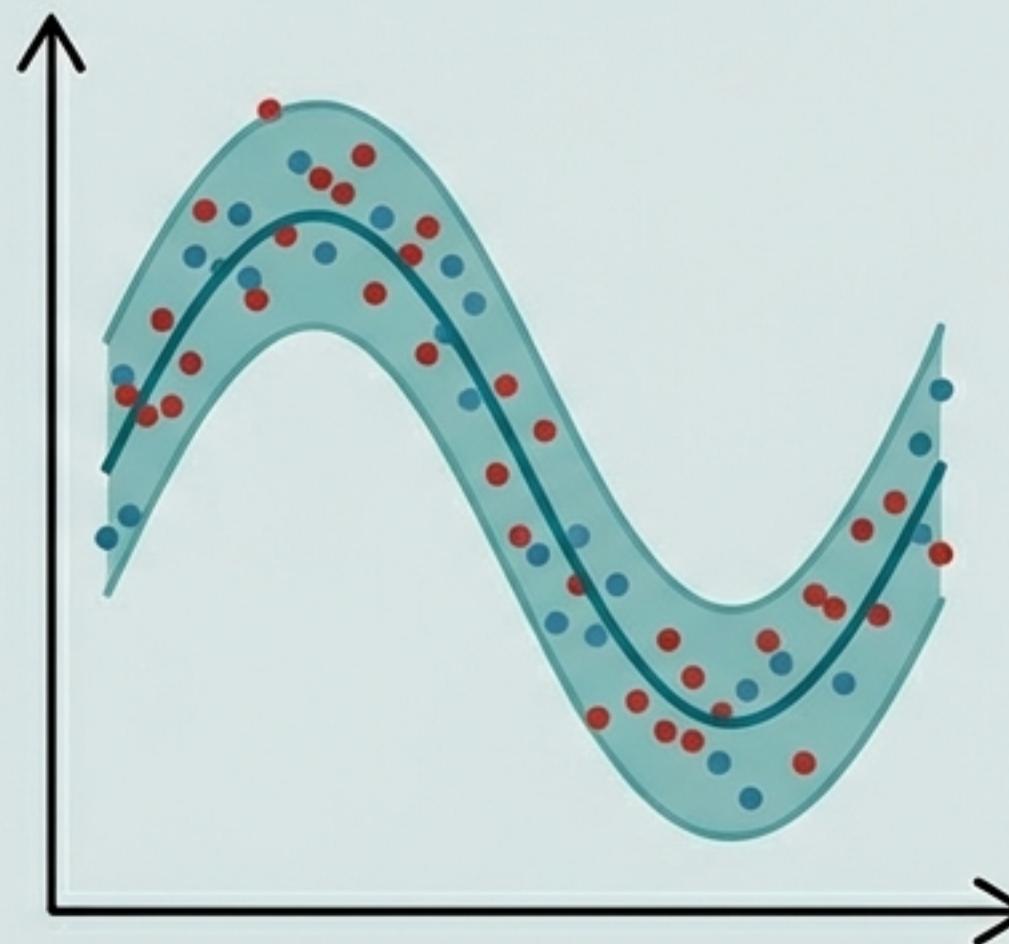
# Tuning Epsilon ( $\epsilon$ ): Setting the Width of the Street

Large  $\epsilon$



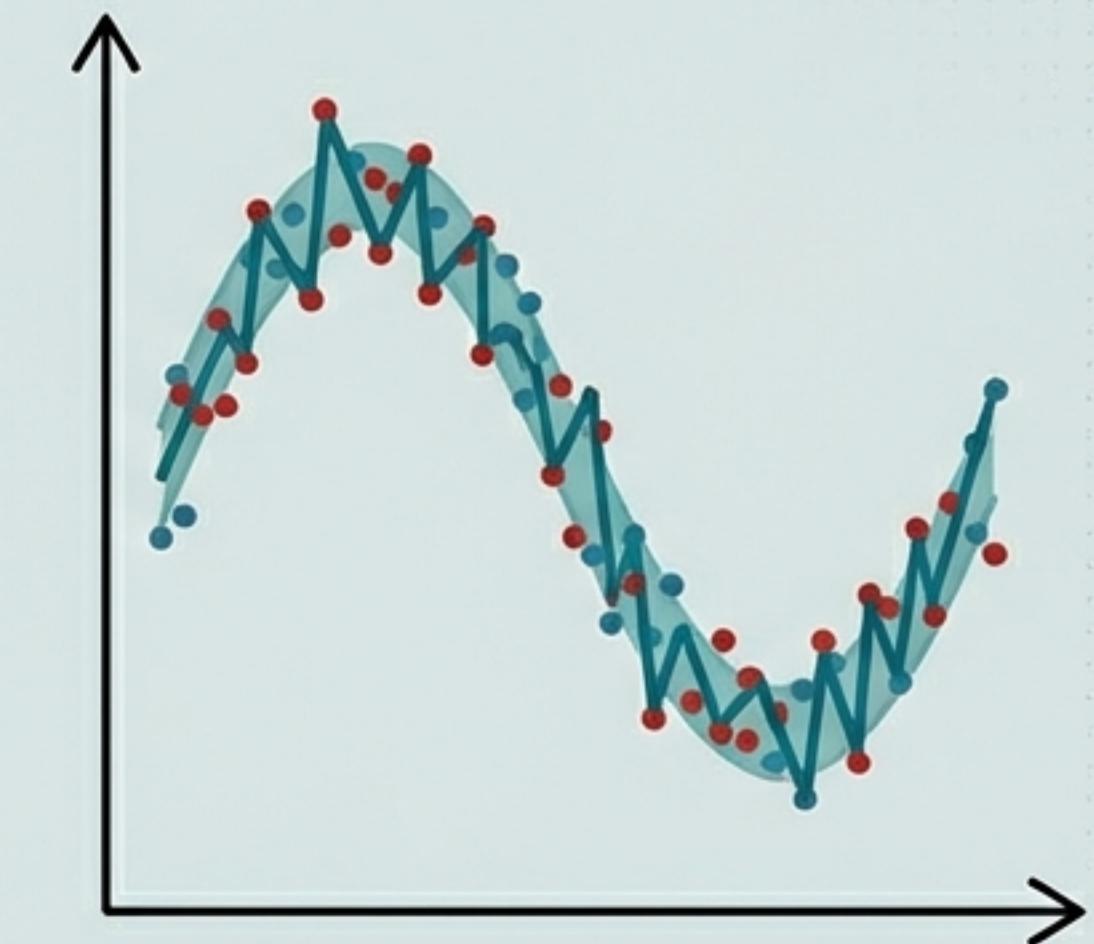
**Underfitting.** High Bias.  
Too tolerant.

Optimal  $\epsilon$



**Balanced.** Captures signal,  
ignores noise.

Small  $\epsilon$

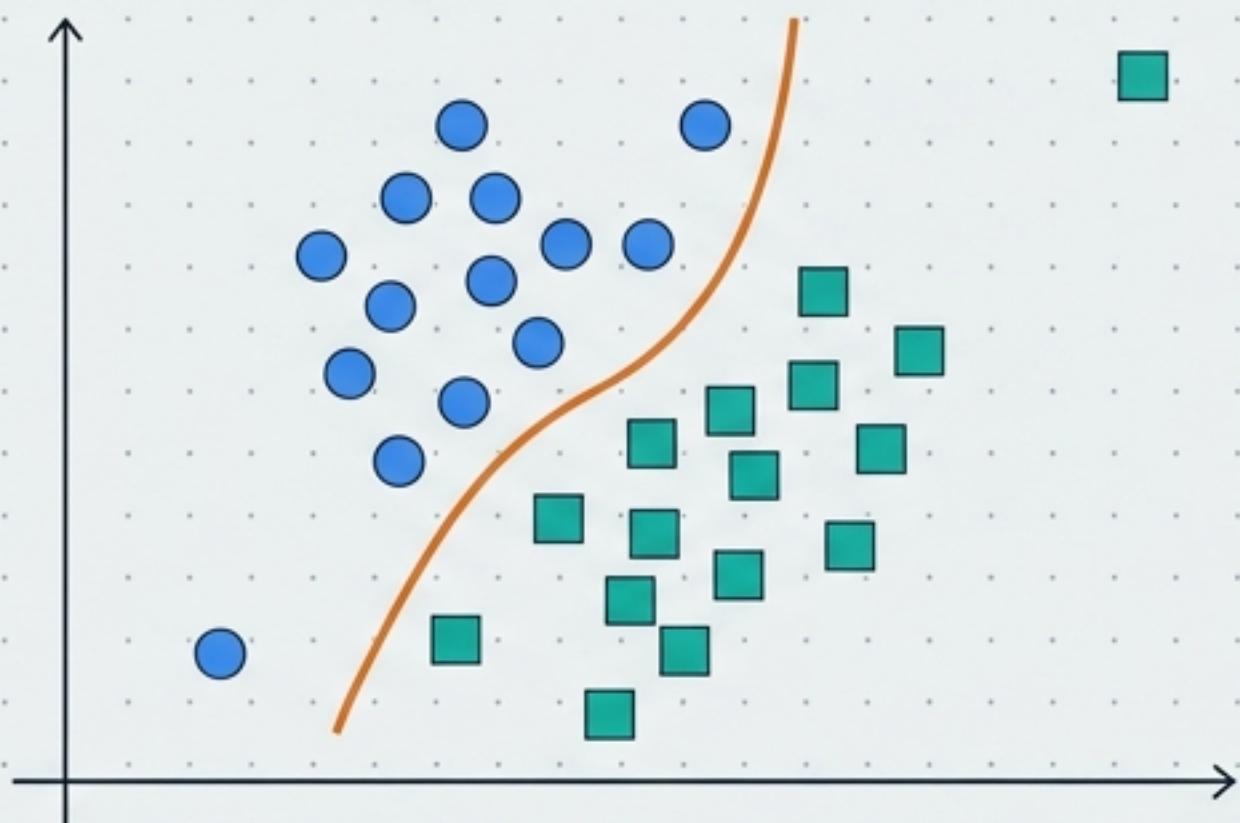


**Overfitting.** High Variance.  
Captures noise.

**Wider Tube = Fewer Support Vectors = Simpler Model.**

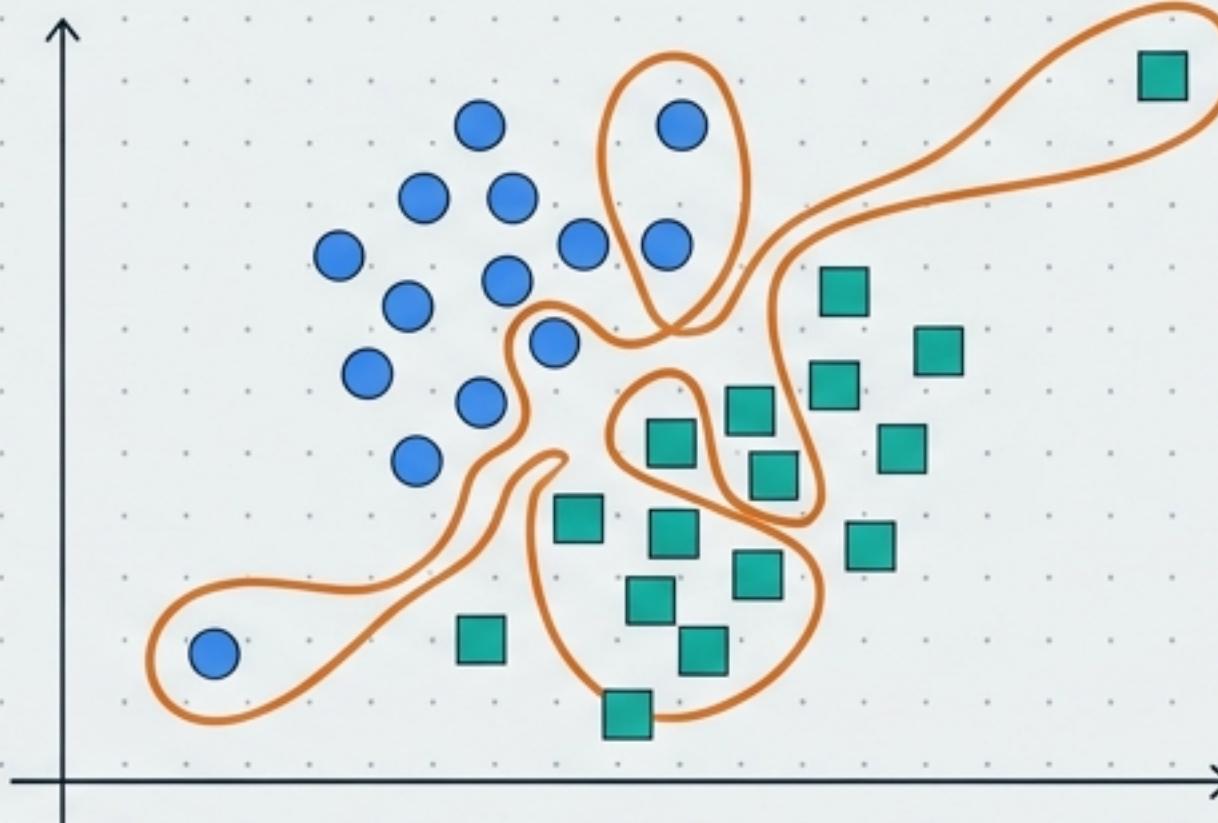
# Tuning Regularization (C): The Penalty for Error

Low C (e.g., 0.1)



**Soft Margin.** "I'm chill, errors are allowed."  
Prioritizes simplicity.

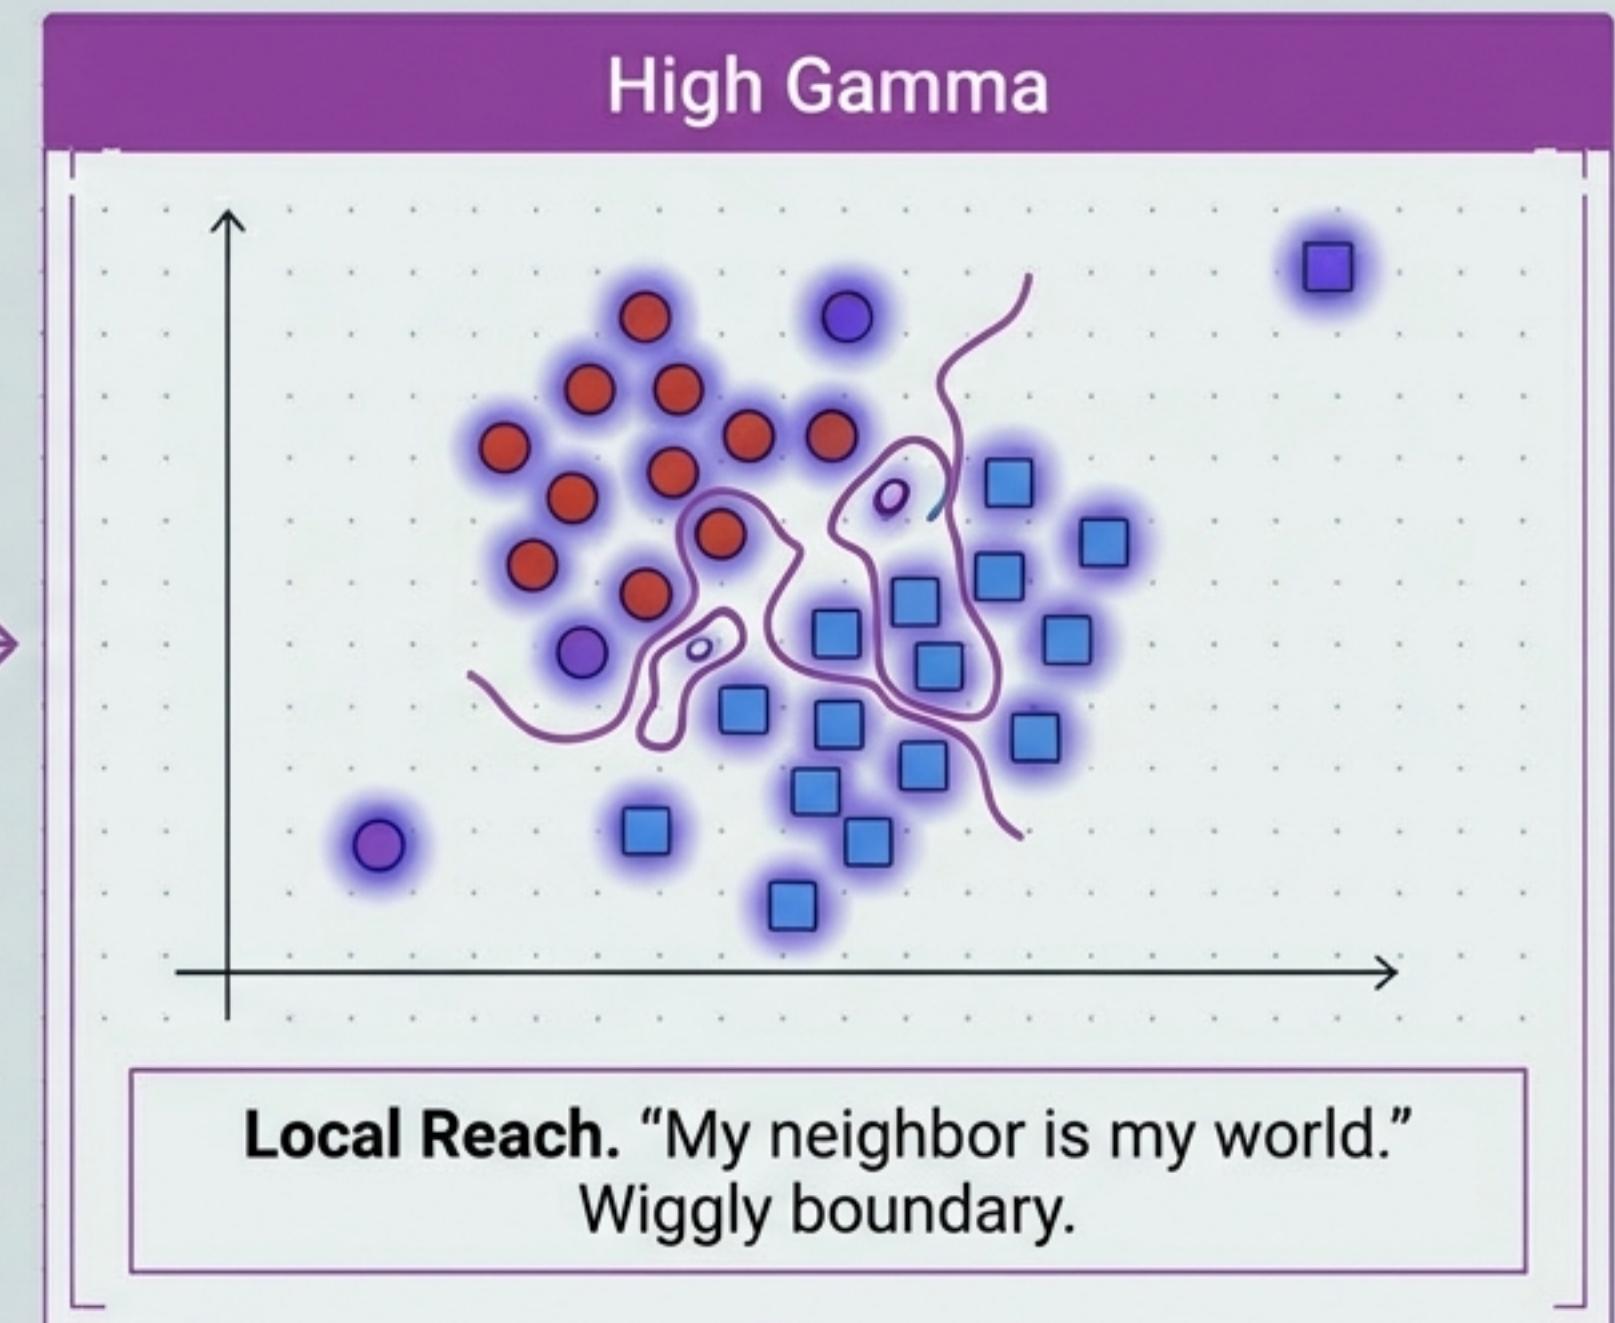
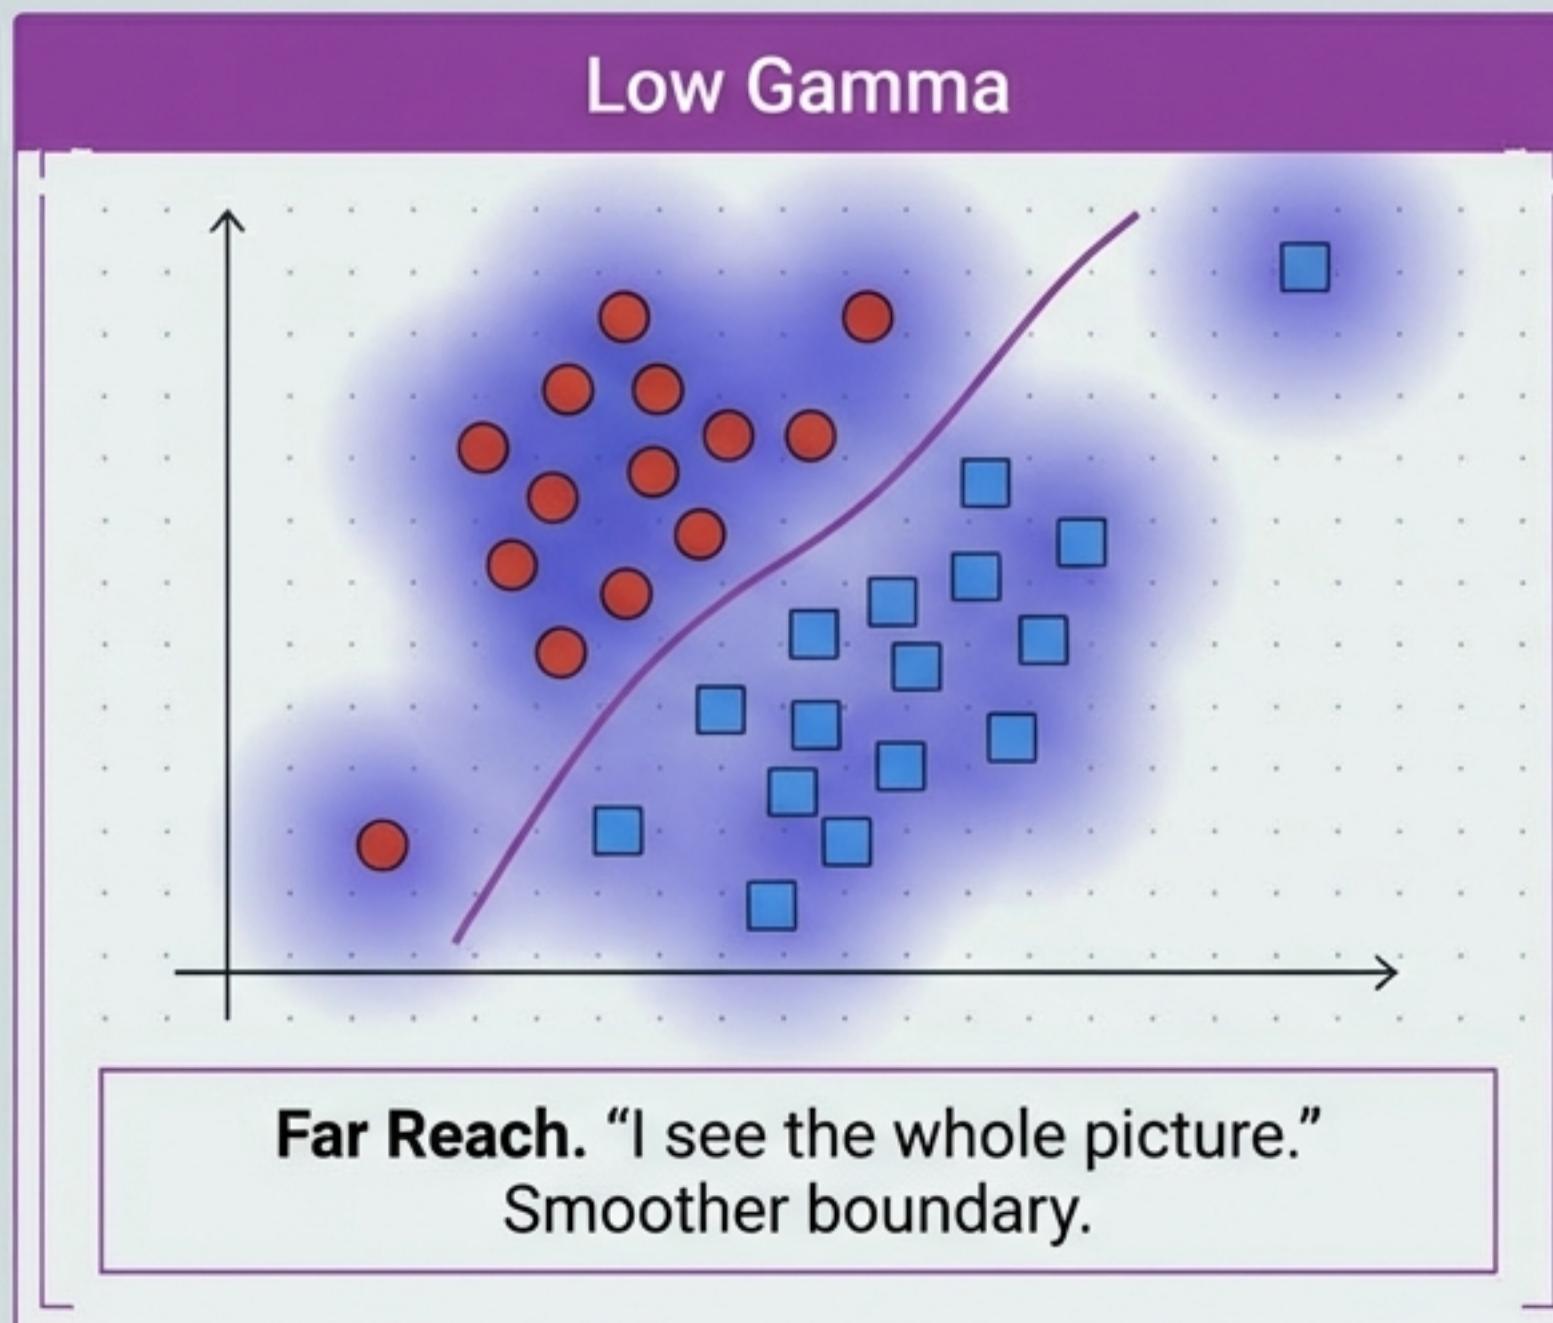
High C (e.g., 1000)



**Hard Margin.** "Strict, zero errors allowed."  
Prioritizes training accuracy (risk of overfitting).

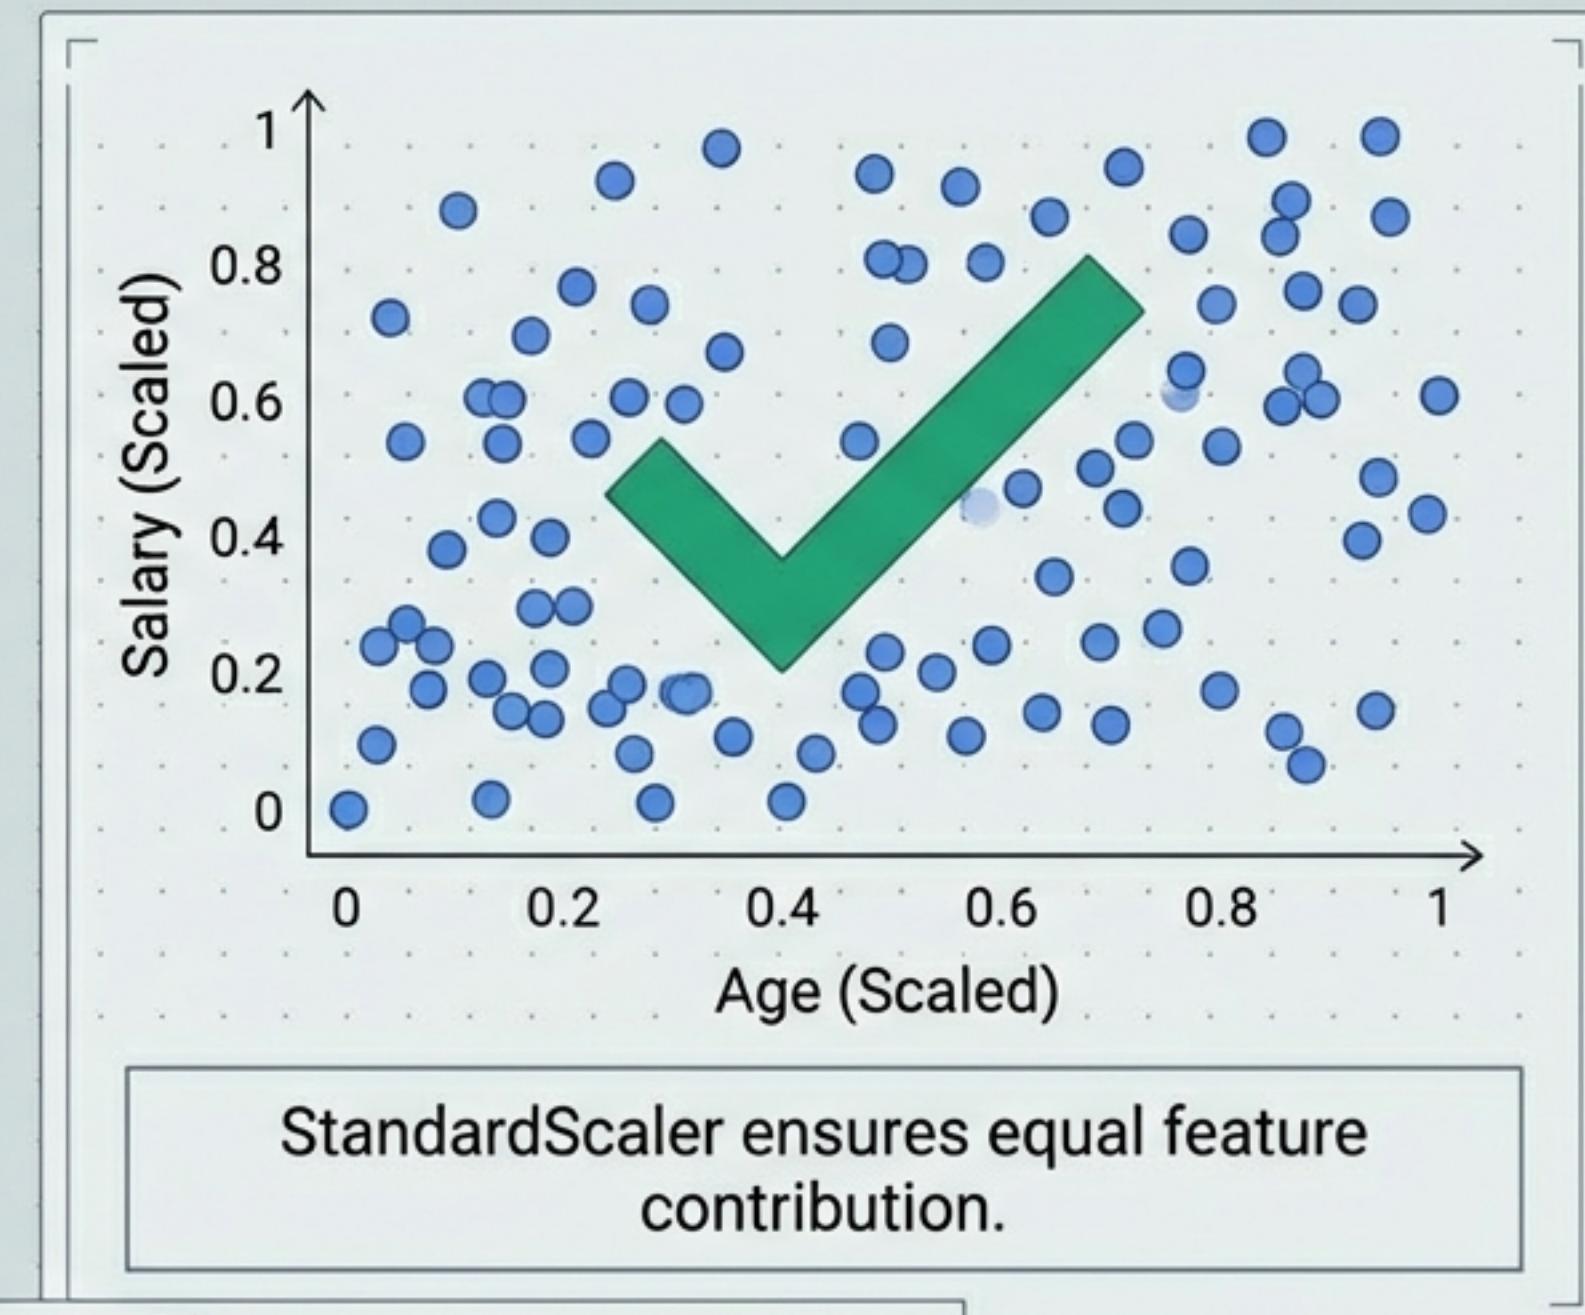
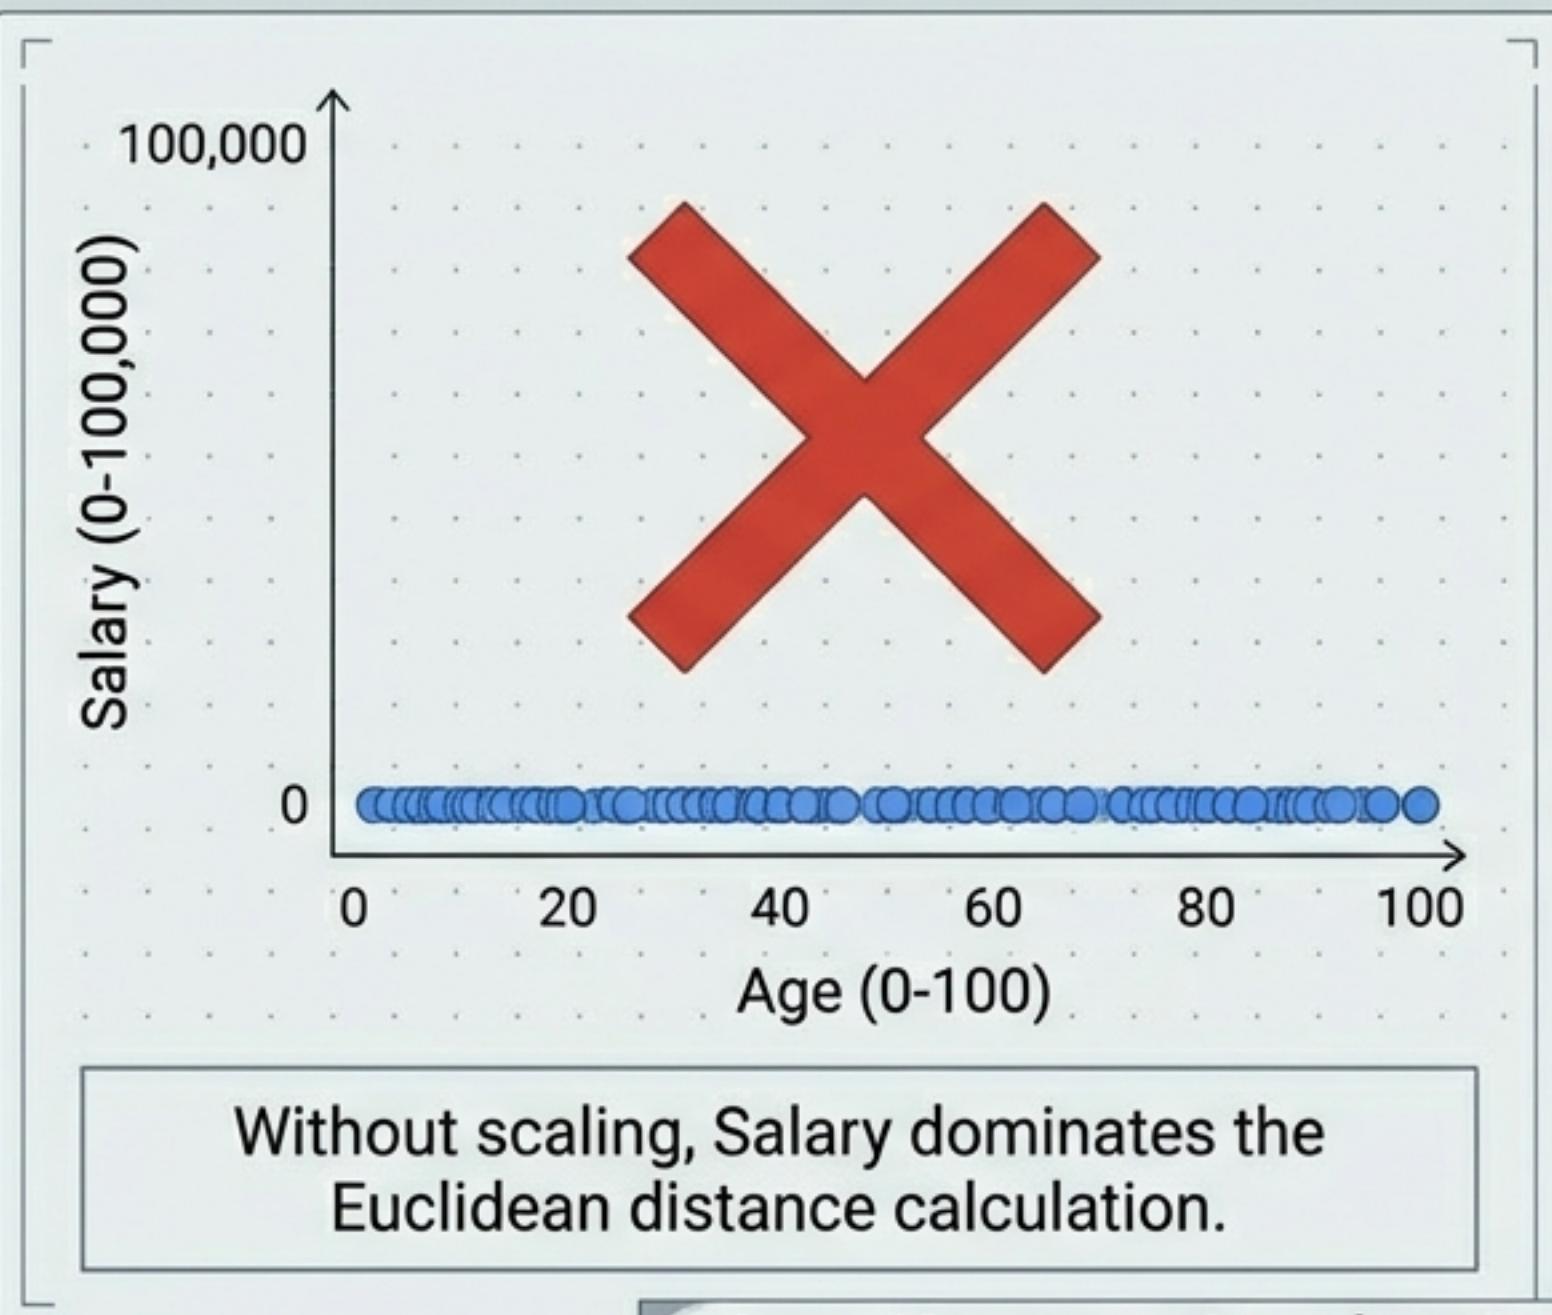
Search Range: Typically Logarithmic ( $10^{-3}$  to  $10^5$ ).

# Tuning Gamma ( $\gamma$ ): The Radius of Influence



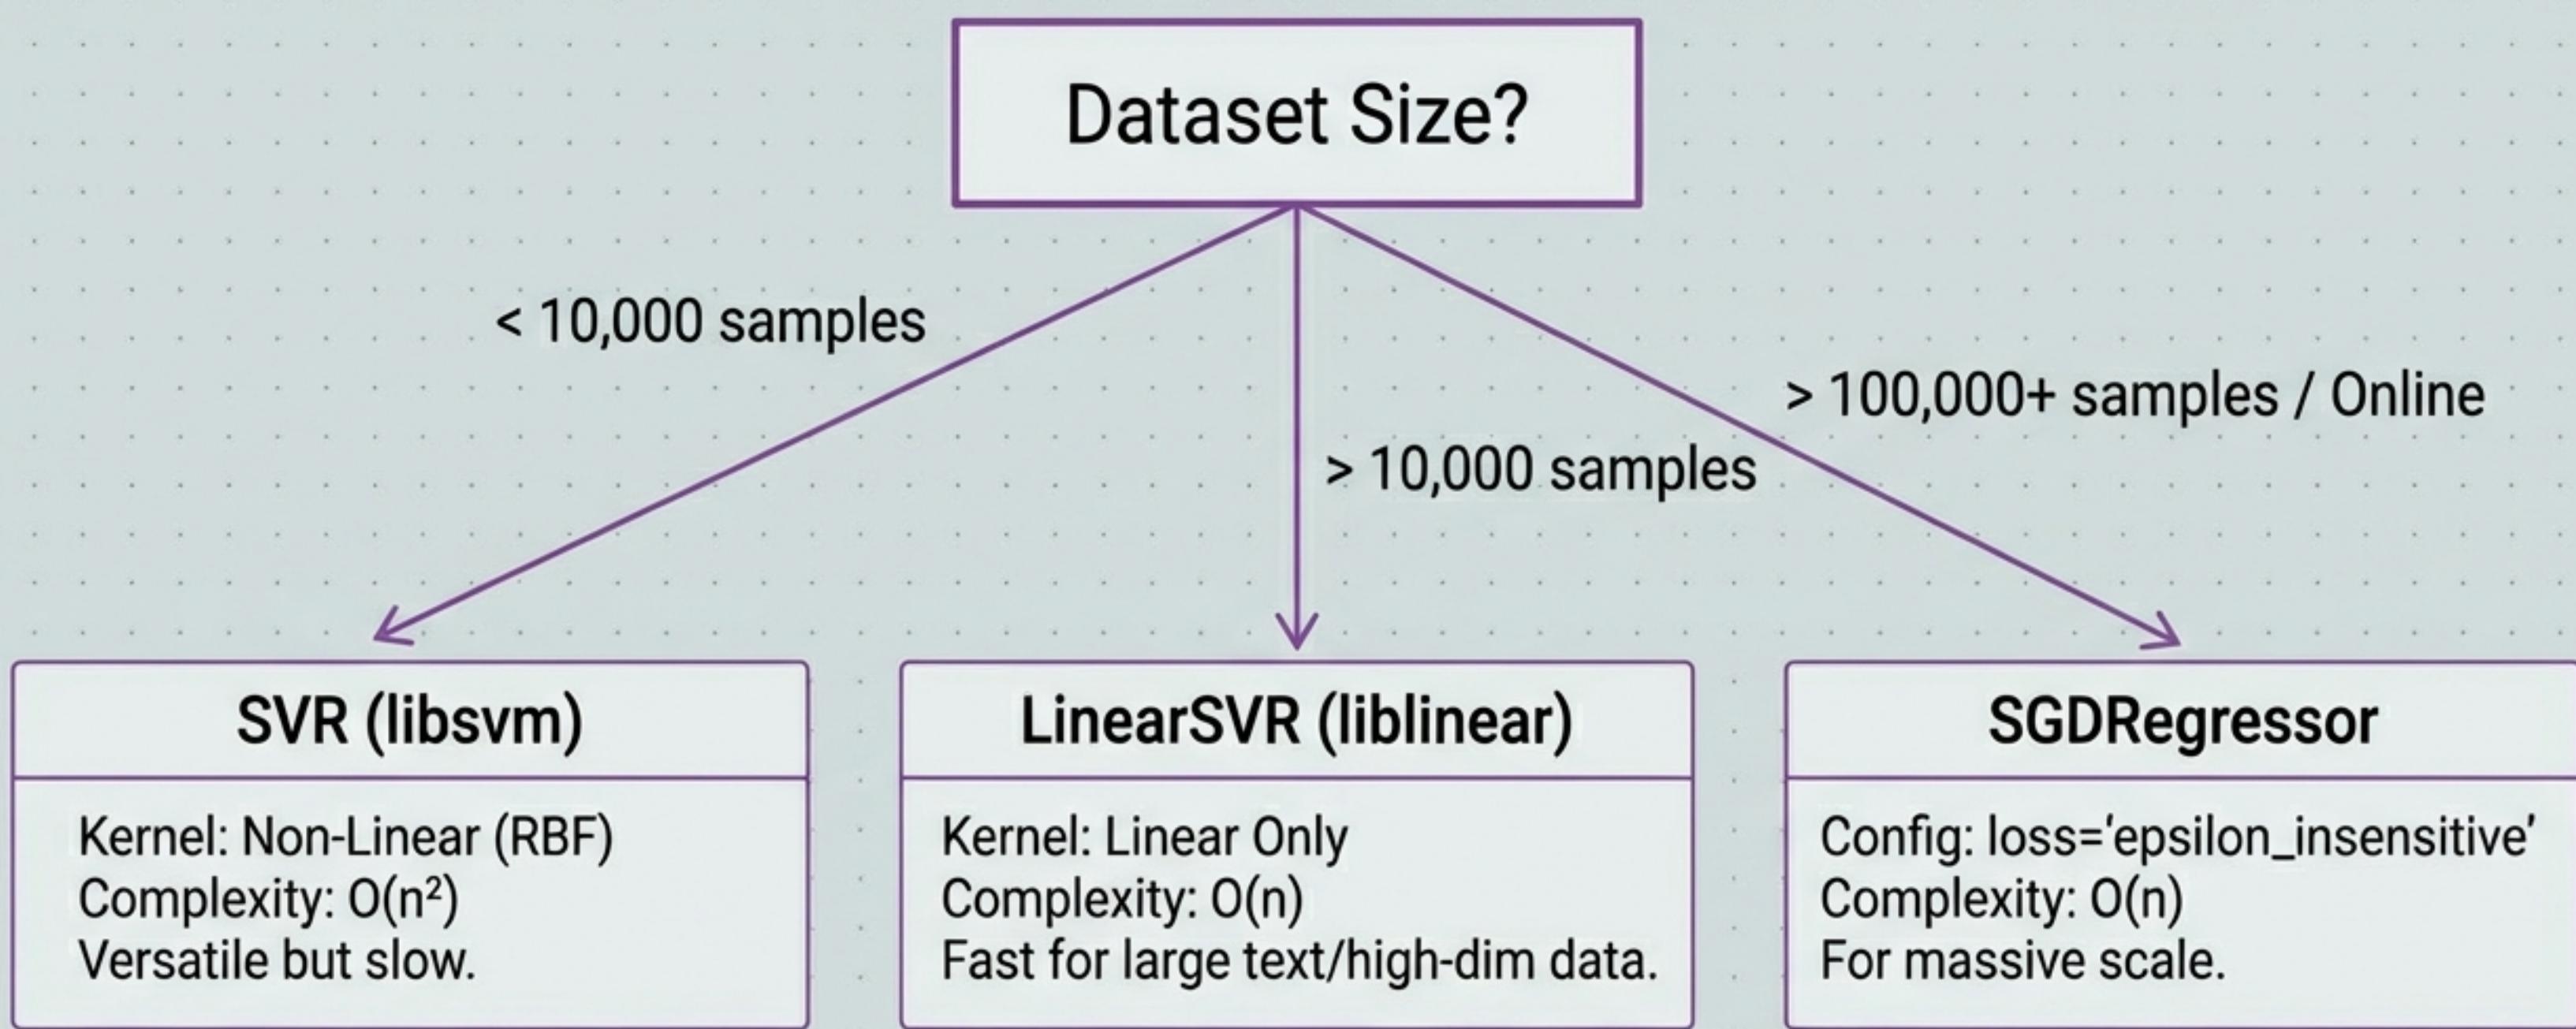
**Default:** `gamma='scale' (1 / (n_features * X.var()))`

# Preprocessing is Not Optional

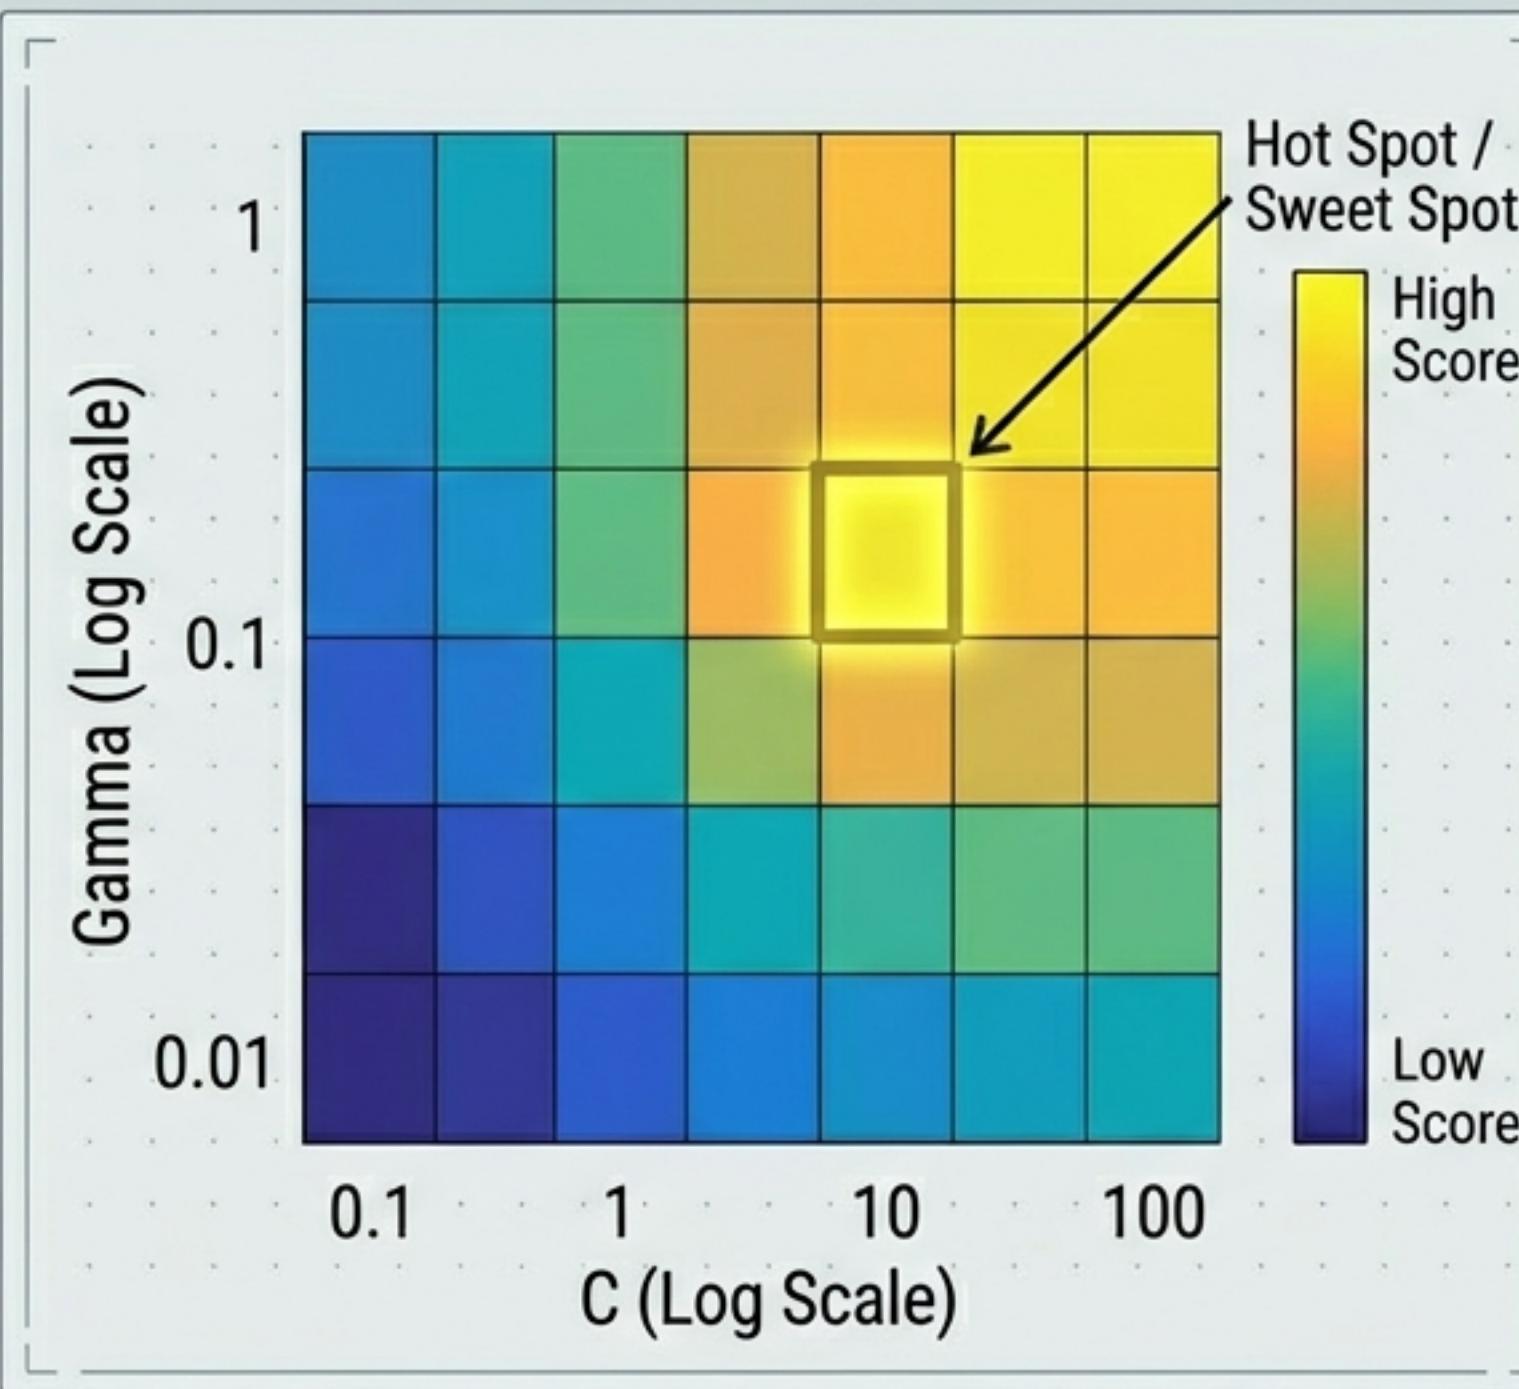


```
from sklearn.pipeline import make_pipeline  
from sklearn.preprocessing import StandardScaler  
from sklearn.svm import SVR  
  
# Always pipeline your scaler!  
regr = make_pipeline(StandardScaler(), SVR(C=1.0, epsilon=0.2))
```

# Scale & Complexity: Choosing the Right Estimator



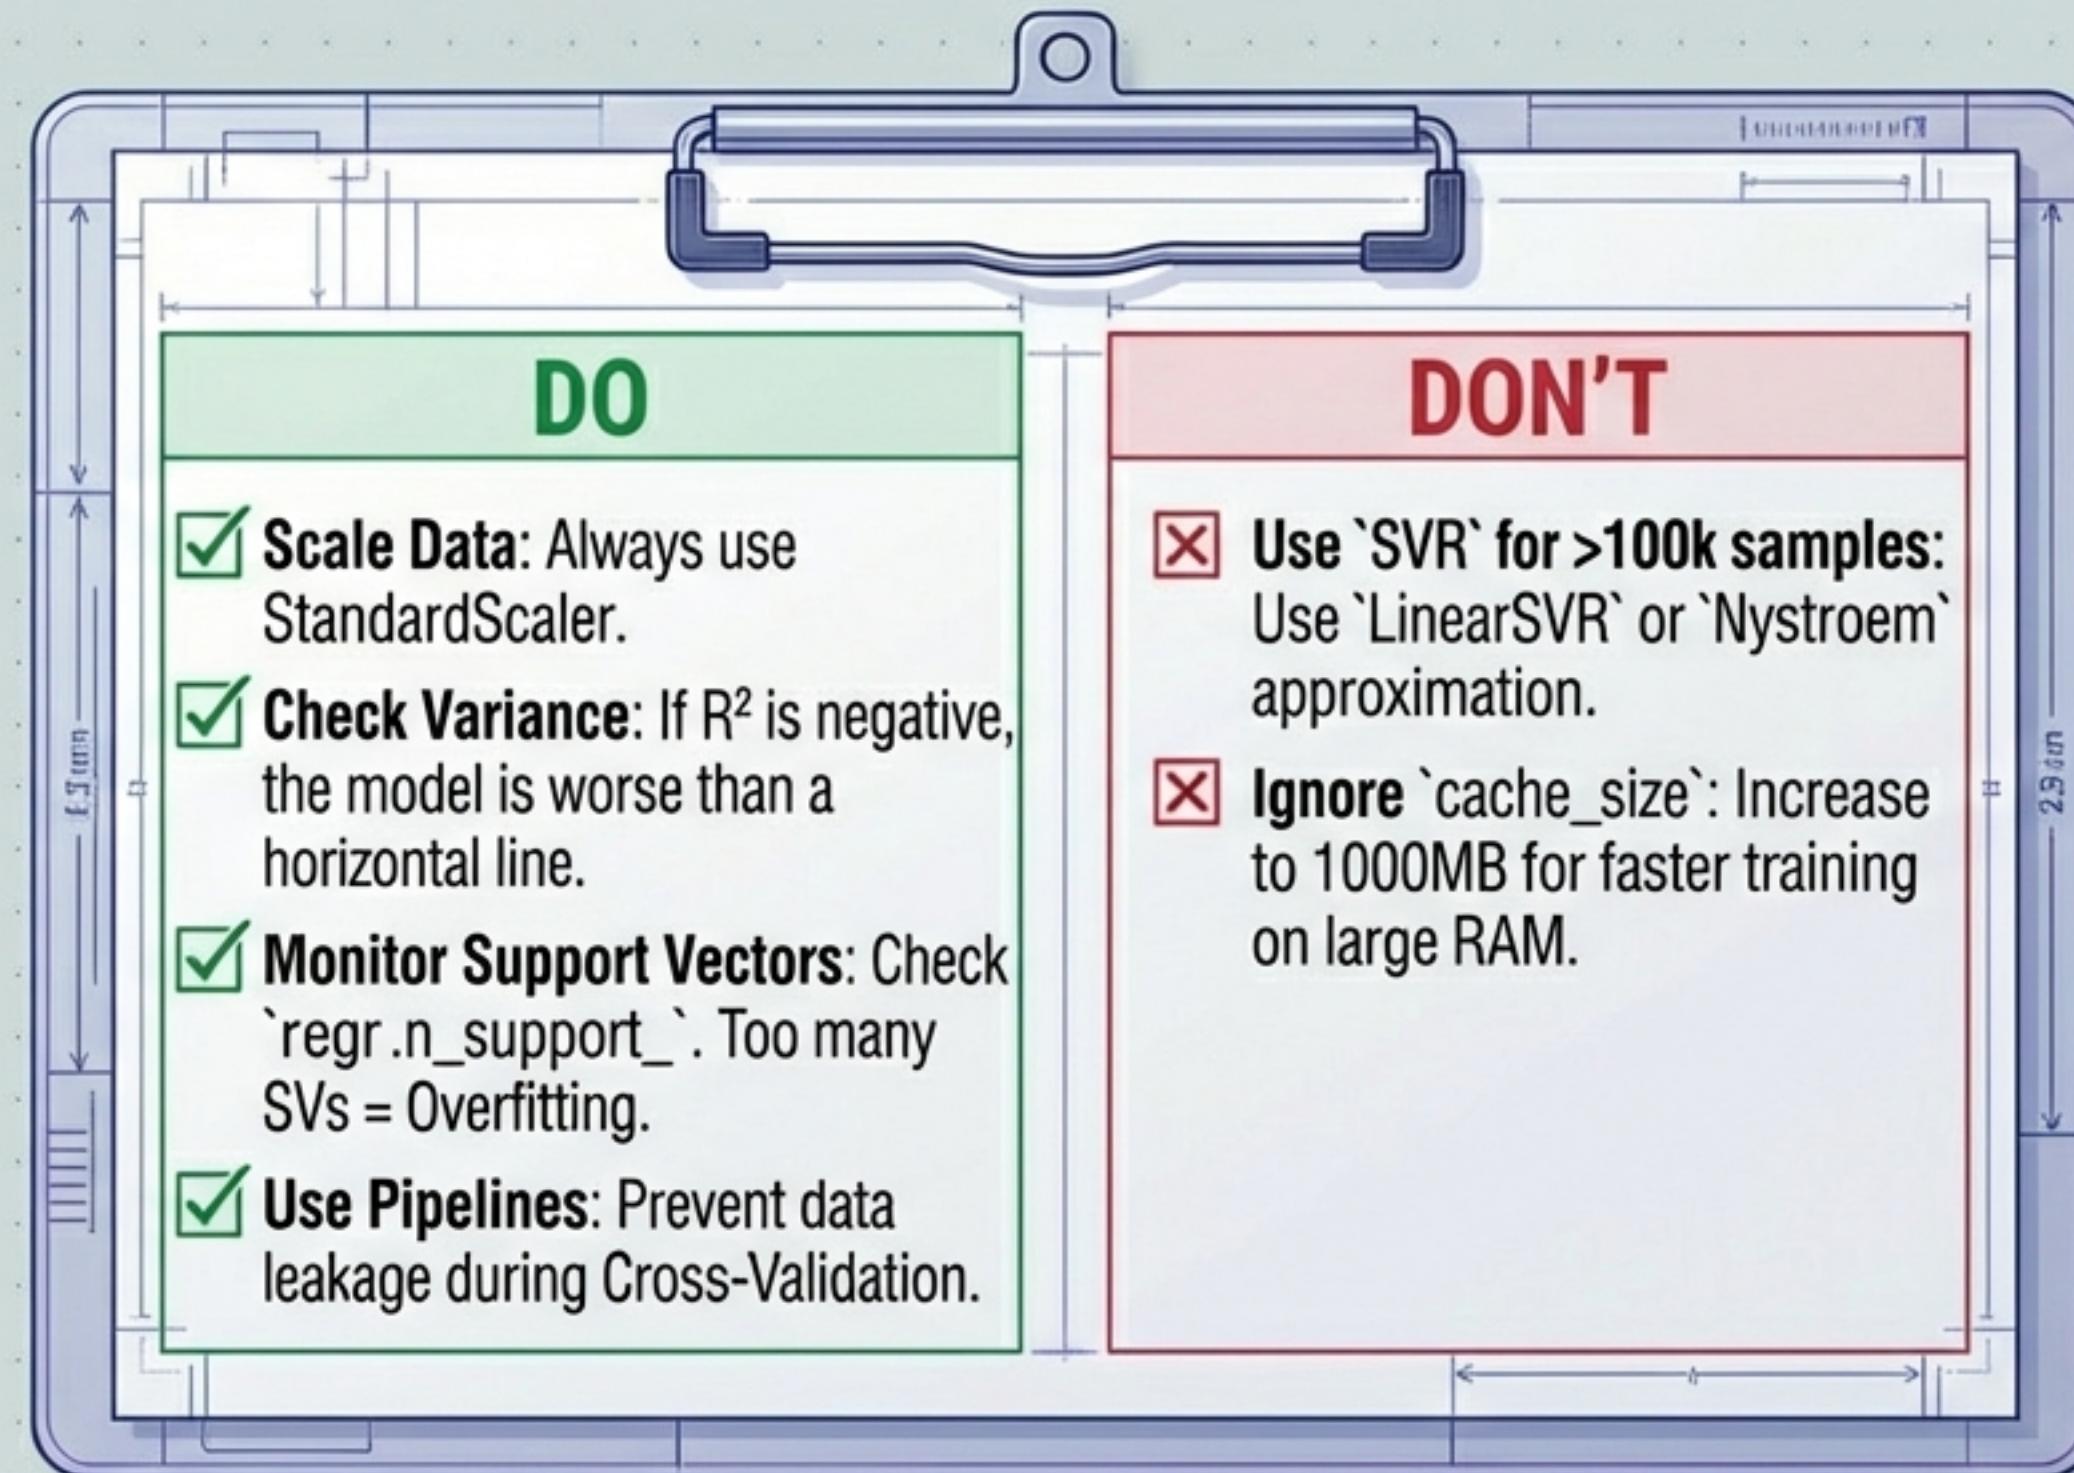
# Search Strategies: Finding the Sweet Spot



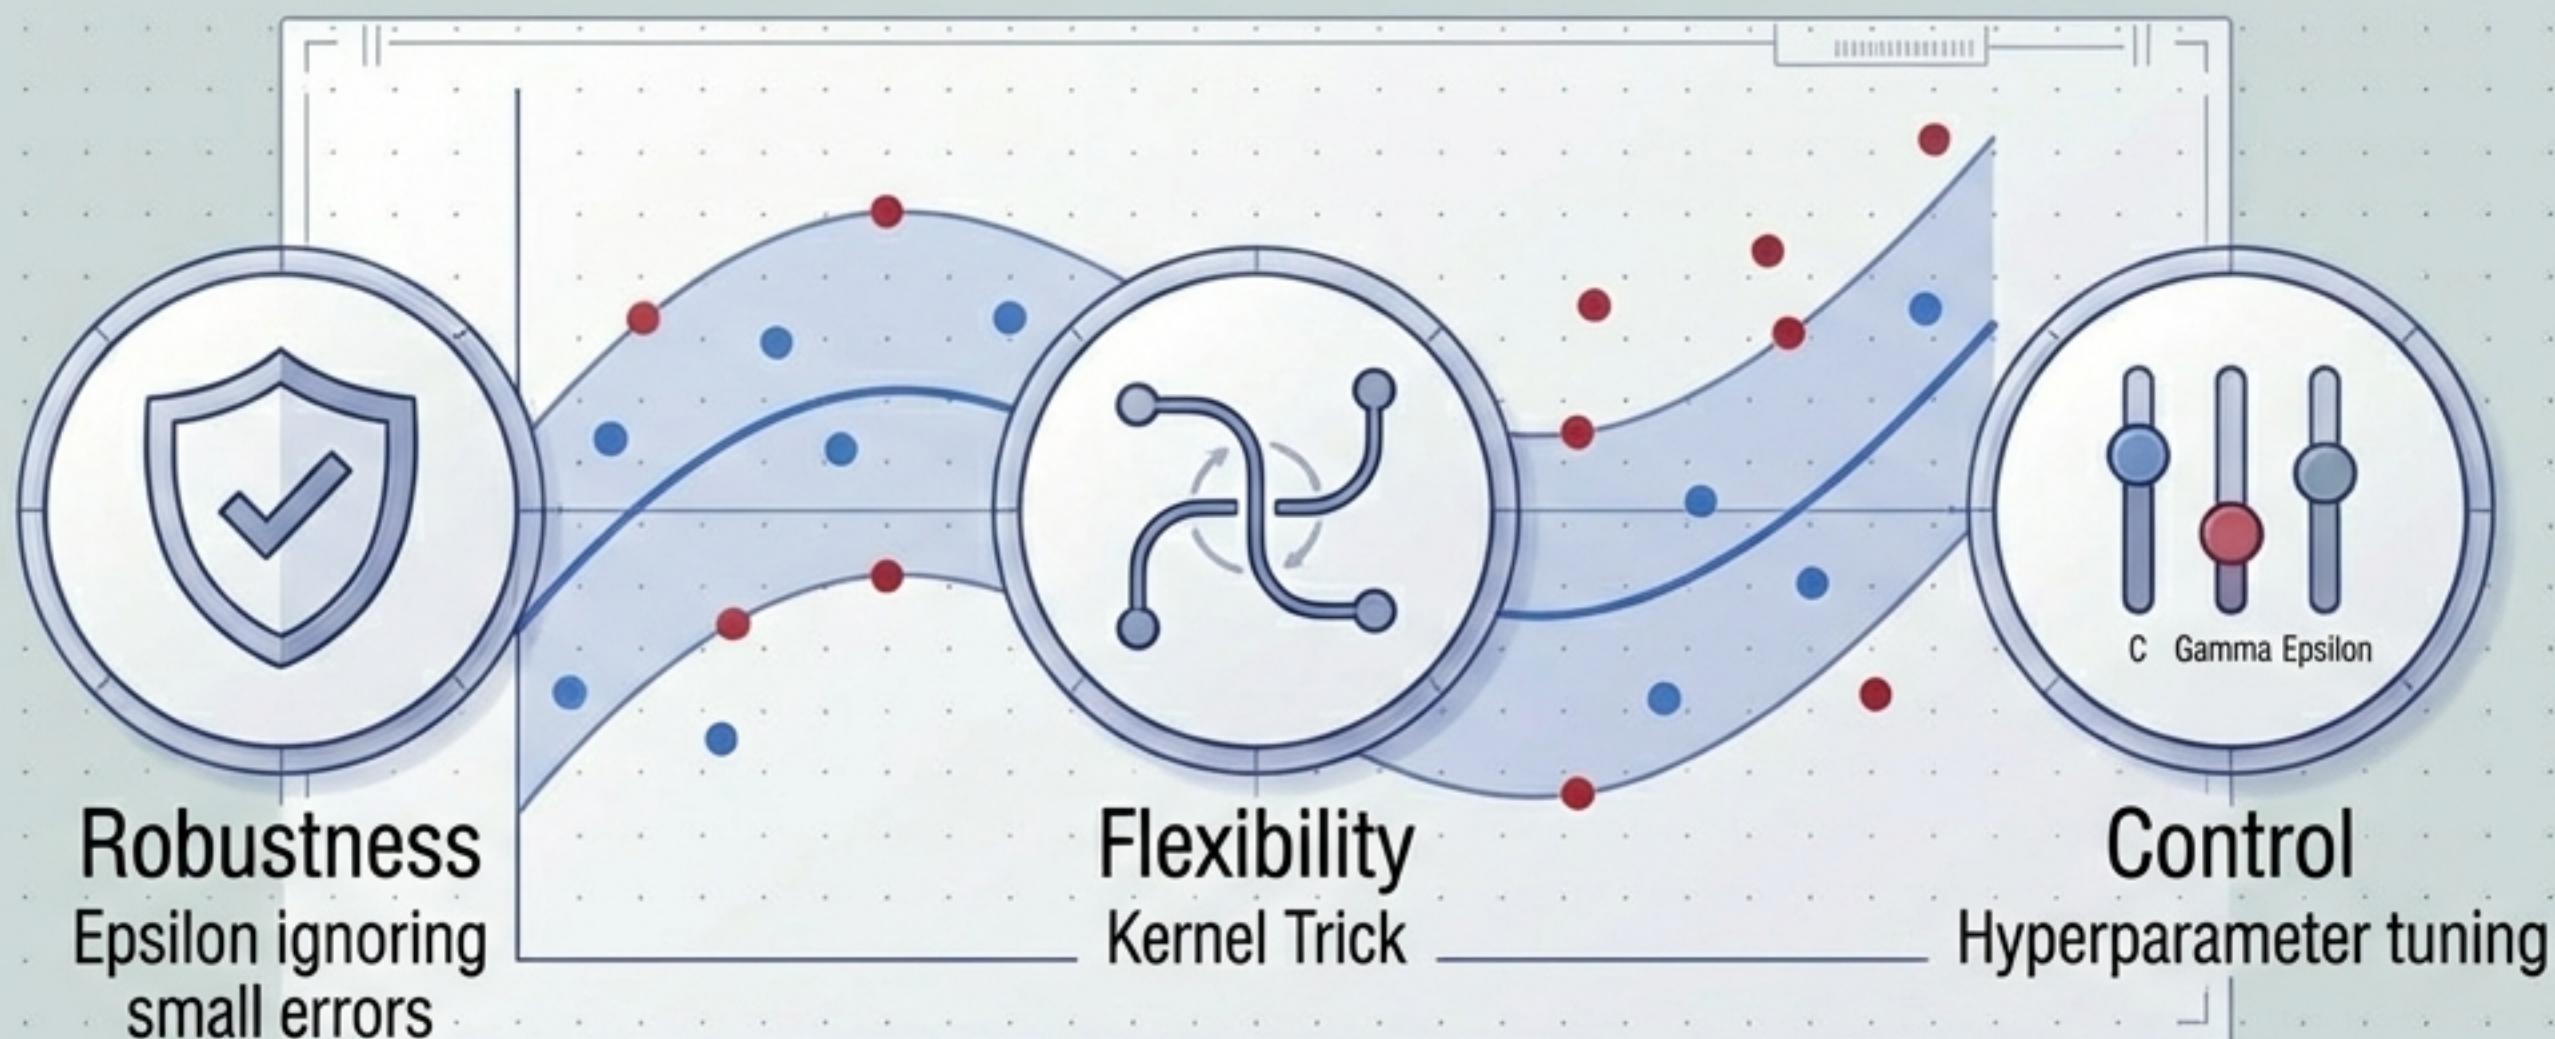
```
param_grid = {  
    'svr__C': [0.1, 1, 10, 100],  
    'svr__epsilon': [0.01, 0.1, 1],  
    'svr__gamma': ['scale', 0.1, 0.01]  
}  
  
grid = GridSearchCV(regr_pipeline,  
param_grid, cv=5)  
grid.fit(X, y)
```

Don't guess. Search.  
Use Logarithmic scales for C and Gamma.

# The Practitioner's Checklist



# Conclusion: The Power of the Tube



SVR is the tool of choice for robust, non-linear modeling when precision matters more than speed.

```
model = SVR(kernel='rbf', C=10, epsilon=0.1)
```